



Integrating Task Parallelism with Actors

OOPSLA, October 25, 2012

Shams Imam, Vivek Sarkar
Rice University



Introduction



- Advent of multi-core processors
- Writing programs exploiting parallelism is hard!
- Renewed interest in parallel and concurrent programming models
- Reduce the burden of reasoning about and writing concurrent programs



Goal



Integrate

- Actor Concurrency (Message Passing)
- Async-Finish (Task Parallelism)

to

- Exploit nondeterministic communication patterns
- Enable easier expression of potential parallelism
- Achieve better performance



Outline

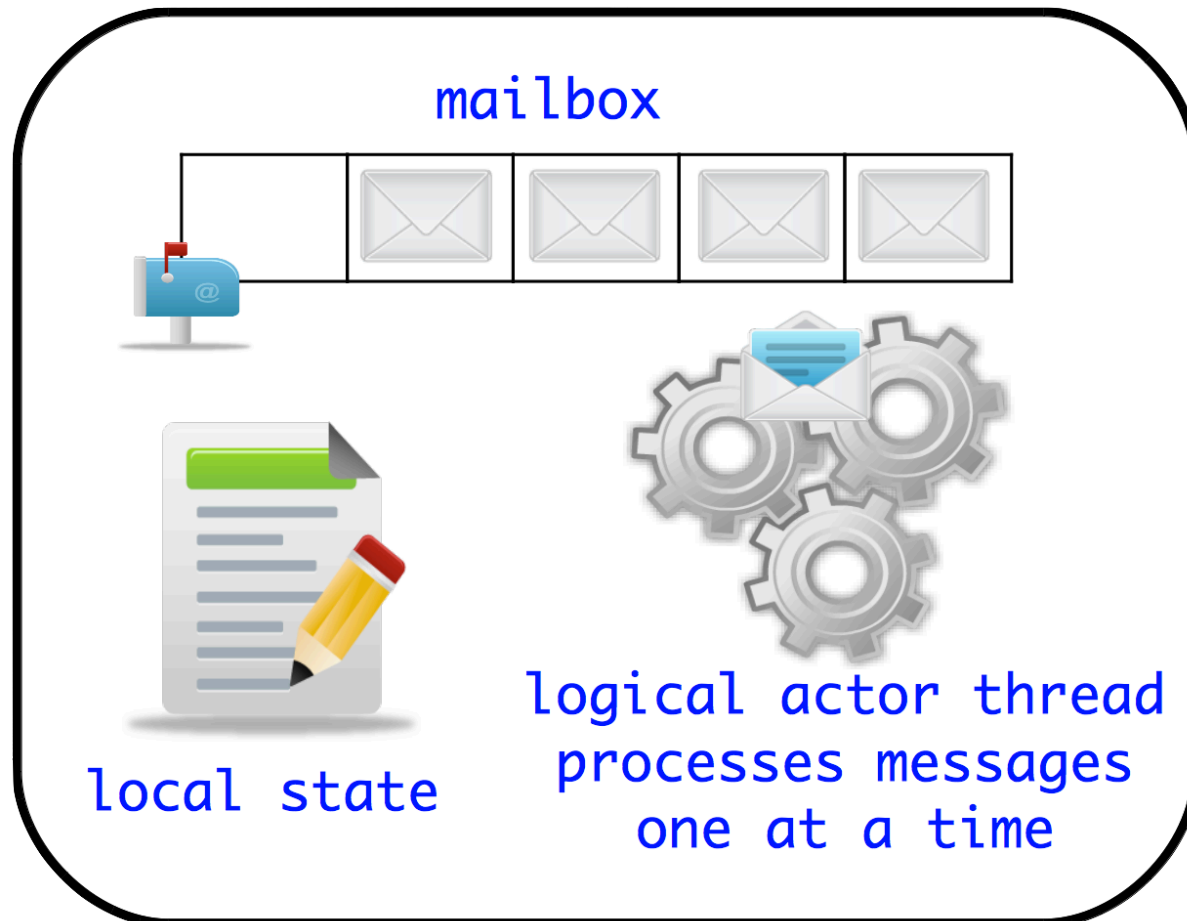


- Introduction
- **The Actor and Async/Finish Models**
- The Unified Model
- Intra-Actor Parallelization
- Experimental Results



The Actor Model

- A message-based concurrency model
- An Actor encapsulates mutable state






Example Scenario



- Pipelined Parallelism
 - each stage can be represented as an actor
 - stages however need to ensure ordering of messages while processing them
 - slowest stage is a throughput bottleneck





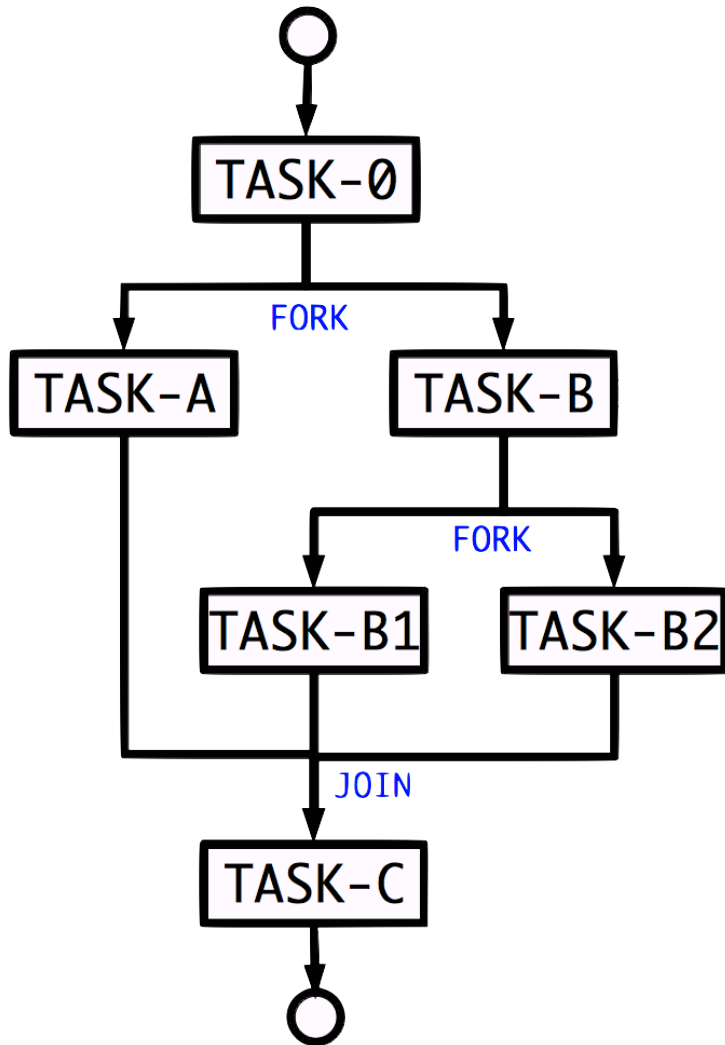
The Async/Finish Model (AFM)



- A special case of the Task Parallel Model
- Parent tasks forks child tasks
- Synchronization when tasks join into another task



AFM Example



```
1  /** Habanero-Scala code **/  
2  object AFMPrimer extends HabaneroApp {  
3    println("Task O"); // Task-O  
4    finish {  
5      async { // Task-A  
6        println("Task A");  
7      }  
8      async { // Task-B  
9        println("Task B");  
10       async { // Task-B1  
11         println("Task B1");  
12       }  
13       async { // Task-B2  
14         println("Task B2");  
15       } } } // Wait for A, B, B1 and B2  
16     println("Task C"); // Task-C  
17   }
```




Outline



- Introduction
- The Actor and Async/Finish Models
- **The Unified Model**
- Intra-Actor Parallelization
- Experimental Results



The Unified Model



- Actors integrated with the AFM
 - Integration needs to be seamless
 - No additional constraints on actors
- Benefits
 - Extend actor capabilities with task parallelism in the unified model
 - Extend task capabilities with more general forms of actor coordination



Actors and Async/Finish Tasks



- Actor creation:
 - instantiate (but do not start) actor instance
 - synchronous operation for creator task (i.e. trivial)
- Actor termination:
 - actor will no longer process messages sent to it
 - synchronous operation by actor (i.e. trivial)
 - all future send requests can be ignored synchronously



Actors and Async/Finish Tasks...




- Starting an Actor:
 - finish scope for actor = finish scope for call to start()
 - actor will start processing messages asynchronously in this finish scope
 - performed by actor task created by runtime
 - actor needs to keep the finish scope “alive” until actor is terminated (even if mailbox is empty)
 - use *lingering* task technique (in a couple of slides)



Actors and Async/Finish Tasks...



- Sending messages:



- possible via *lingering* task technique



Lingering Tasks



- Provide a hook into some finish scope
- Use the *lingering* task to spawn new send and message processing tasks
- One *lingering* task per actor
 - created when the actor is started
 - *lingering* task completes execution only when the actor terminates



Easier Termination Detection in Unified Model



```
1 /** Scala code **/  
2 object Terminator extends App {  
3   val latch = new CountdownLatch(1)  
4   val printActor = new PrintActor(latch)  
5   printActor.start()  
6   printActor.send("Hello World")  
7   printActor.send(StopMsg())  
8   latch.await()  
9   println("Actor terminated")  
10 }  
  
12 class PrintActor(latch: CountdownLatch)  
13   extends Actor {  
14   def act() {  
15     loop { react {  
16       case msg: String =>  
17         println(msg)  
18       case msg: StopMsg =>  
19         latch.countDown()  
20       exit()  
21     } } } }
```

```
1 /** Habanero-Scala code **/  
2 object Terminator extends HabaneroApp {  
3   val printActor = new PrintActor()  
4   finish {  
5     printActor.start()  
6     printActor.send("Hello World")  
7     printActor.send(StopMsg())  
8   } // wait until actor terminates  
9   println("Actor terminated")  
10 }  
  
12 class PrintActor extends UnifiedActor {  
13   def behavior() = {  
14     case msg: String =>  
15       println(msg)  
16     case msg: StopMsg =>  
17       exit()  
18   } }
```



Outline



- Introduction
- The Actor and Async/Finish Model
- The Unified Model
- **Intra-Actor Parallelization**
- Experimental Results



Parallelizing Actors



- Traditionally actor message processing (MP) has been sequential
- Under the AFM, we can use two techniques to parallelize the MP...



Parallelizing Actors (contd)



1. Use finish construct in MP body and spawn child tasks


```
1  /** Habanero-Scala code */
2  class FirFilter(..., nextStage: UnifiedActor)
3      extends UnifiedActor {
4      ...
5      def behavior() = {
6          case FirItemMessage(value, coeffs) =>
7              ...
8              val helpers = ... // number of helper tasks
9              val stores = Array.ofDim[Double](helpers)
10             finish {
11                 // compute the sum using divide-and-conquer
12                 (0 until helpers) foreach { helperId =>
13                     async {
14                         val (start, end) = ...
15                         var sum: Double = 0.0
16                         start until end foreach { index =>
17                             sum += buffer(index) * coeffs(index)
18                         }
19                         stores(helperId) = sum
20                     } } }
21             // propagate the sum down the pipeline
22             val globalSum = stores.foldLeft(0.0) {
23                 (acc, loopVal) => acc + loopVal
24             }
25             nextStage.send(DataItemMessage(globalSum))
26             case ... => ...
27         } }
```



Example Scenario



- Pipelined Parallelism
 - reduce effects of slowest stage by introducing task parallelism
 - increases the throughput





Parallelizing Actors (contd)



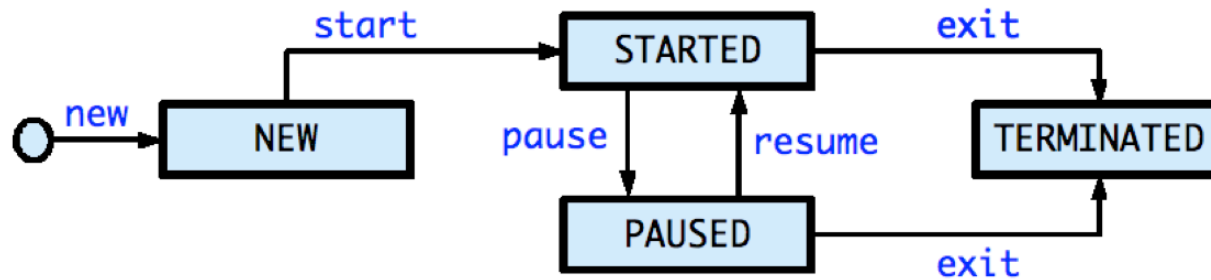
2. Allow *escaping* asyncs inside MP body

```
1  /** Habanero—Scala code */
2  class ParallelizedActor() extends UnifiedActor {
3      override def behavior() = {
4          case msg: SomeMessage =>
5              finish {
6                  async { ... /* processing in parallel */ }
7                  // some more processing
8              }
9          async { ... /* escaping async */ }
10     ...
11 } }
```

- **WAIT!** What about the single message processing invariant?



Pause and Resume an Actor



- paused state
 - actor will no longer process messages sent to it
- new operations:
 - `pause()`: move from started to paused state
 - `resume()`: move from paused to started state
- pause actor before returning from MP body
- resume actor when safe to process next message



Parallelizing Actors (contd)



```
1  /** Habanero-Scala code */
2  class EscapingAsyncActor() extends UnifiedActor {
3    override def behavior() = {
4      case msg: SomeMessage =>
5        async { /* do some processing in parallel */ }
6        // preprocess the message
7        pause() // delay processing the next message
8        // pause/resume is not thread blocking
9        async {
10          // do some more processing in parallel
11          // safe to resume processing other messages
12          resume()
13          // some more processing
14        } ...
15    } }
```



Outline



- Introduction
- The Actor and Async/Finish Model
- The Unified Model
- Intra-Actor Parallelization
- **Experimental Results**



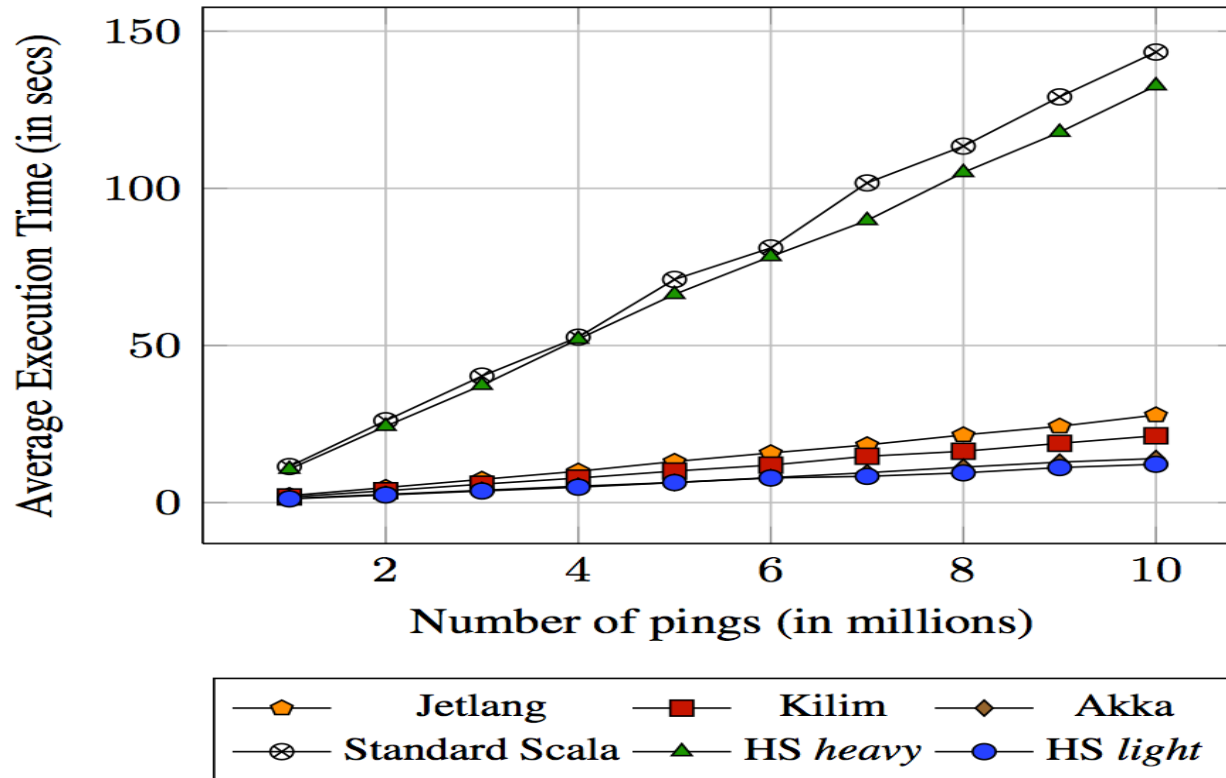
Experimental Setup



- Configuration
 - 2.8 GHz Intel Westmere, 12-core SMP node
 - 48 GB of RAM per node (4 GB per core)
 - Red Hat Linux (RHEL 6.0)
 - Sun Hotspot JDK 1.7
 - Scala 2.9.1-1
 - Habanero-Scala 0.1.3 (<http://habanero-scala.rice.edu/>)
- Execution time reported using variant of “Statistically Rigorous Java Performance Evaluation” by A. Georges et al.



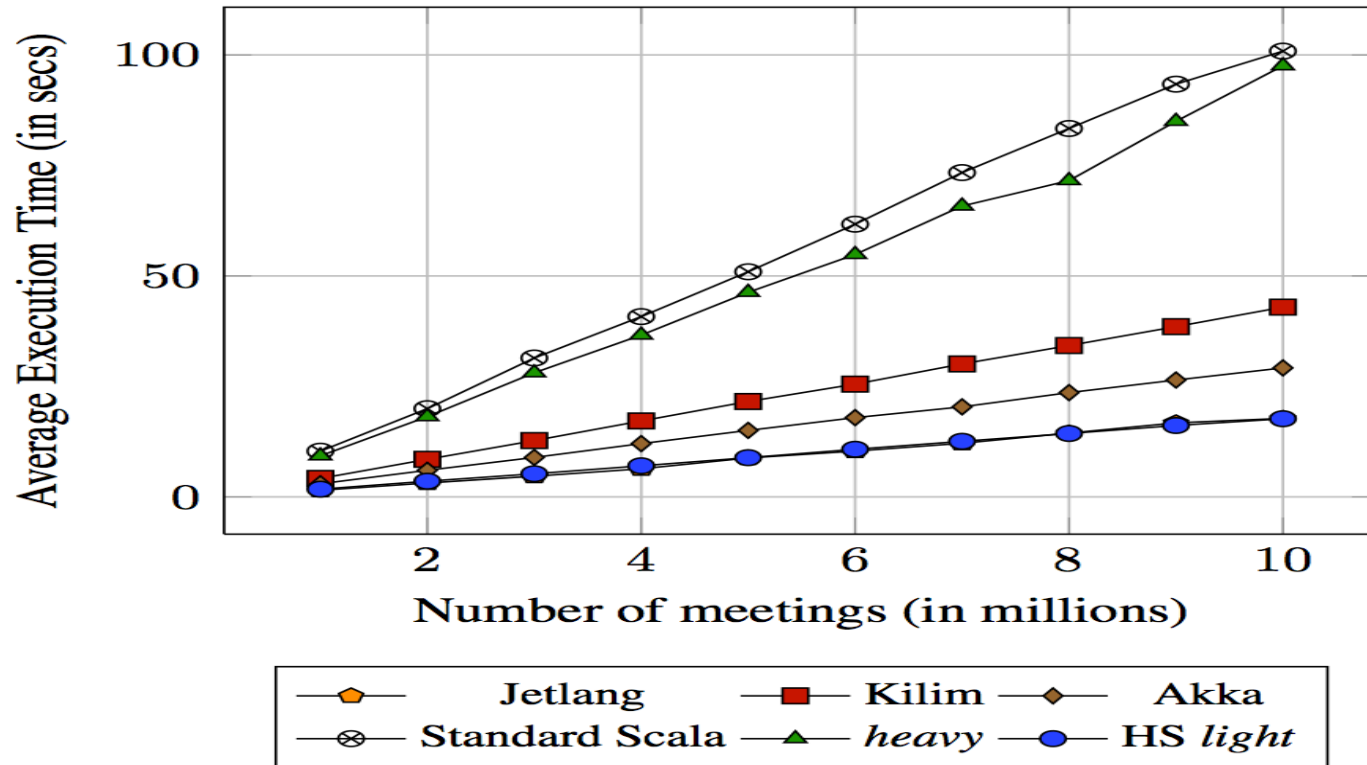
Ping-Pong Benchmark



- measures raw message throughput
- HS Light and Akka actors fastest
 - no exceptions
 - Fork-join scheduler



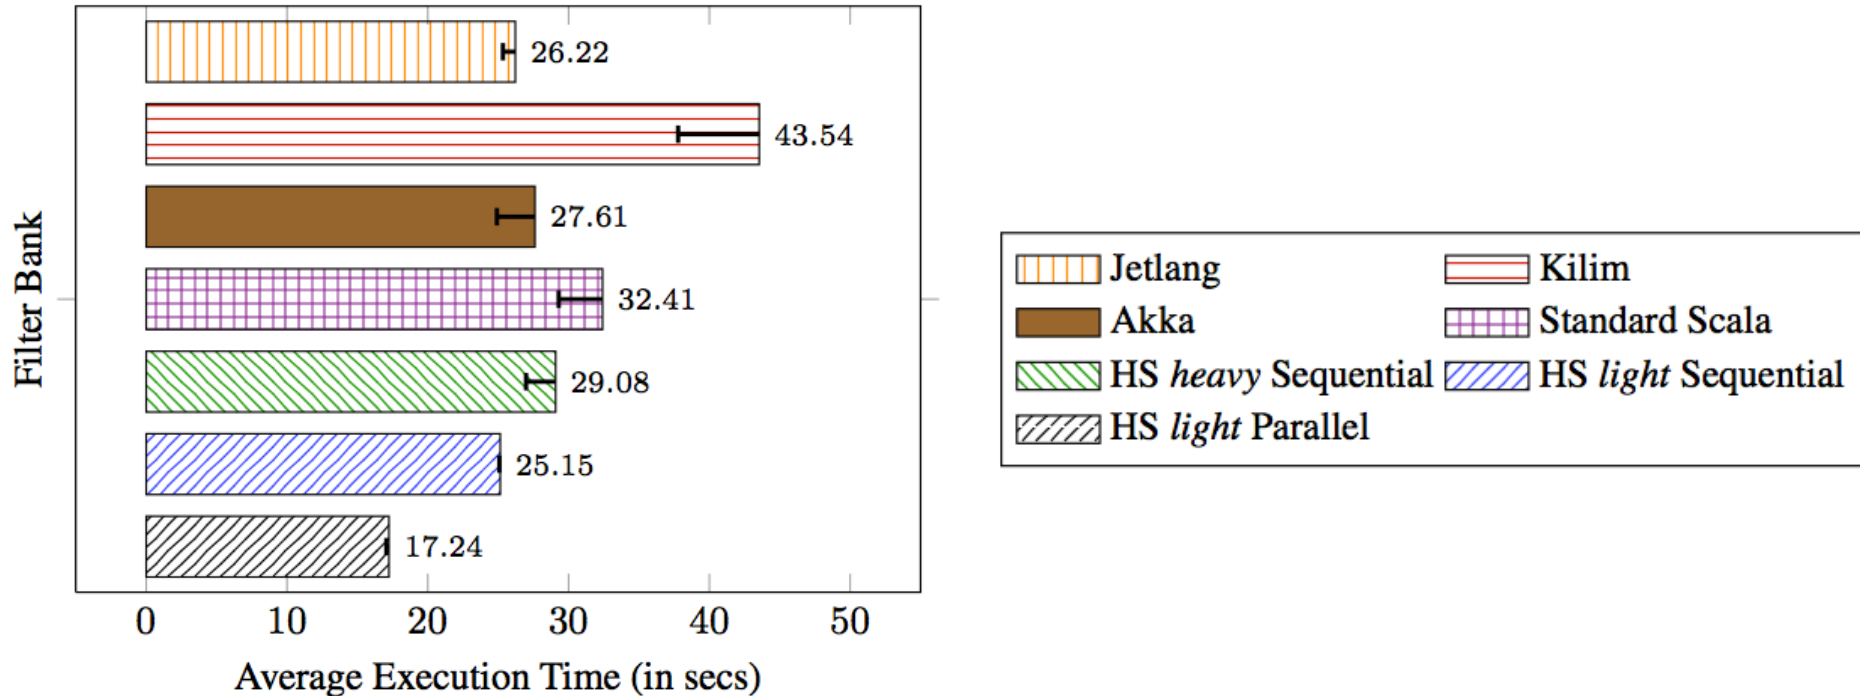
Chameneos Benchmark



- measures cost of synchronization in the mailbox
- HS Light actor and Jetlang fastest
 - both uses batch processing of messages
 - Light actors use DDCs



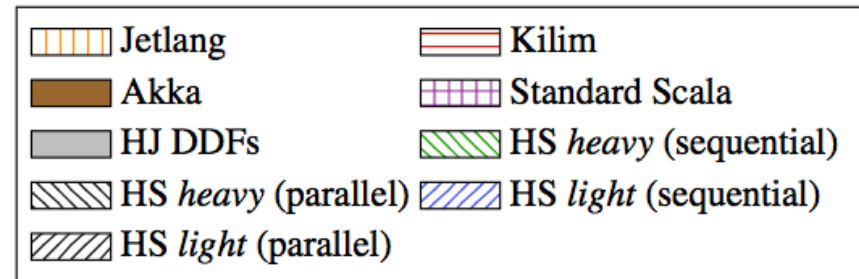

Filter Bank (Pipeline)



- Unified solution fastest
 - up to 30% faster than pure Actor solutions
 - stage parallelization shortens critical length of the pipeline



Quicksort Benchmark



- Unified solution fastest
 - up to 10% faster than pure Actor solutions
 - around 23% faster than DDF solution



Related Work



- Parallel Actor Monitors by Scholliers et al.
 - Parallelism by modifying message processing scheduler
 - Does not allow parallelism inside message processing body
- CoBox Model by Schäfer et al.
 - Communicating CoBoxes host multiple objects
 - Each CoBox contains multiple tasks but cooperatively executes only one at a time
 - Our model is equivalent to allowing a CoBox to execute multiple tasks at a time
- Other related work discussed in the paper



Summary



- A unified programming model that integrates
 - the Async/Finish model
 - the Actor model
- Application characteristics that benefit from the unified model
- Future work
 - Extend past work on data race detection for the AFM to the unified model



Acknowledgments



- U.S. NSF awards 0926127 and 0964520
- Habanero Group
 - Vincent Cavé
 - Dragoş Sbîrlea
 - Saĝnak Taşırlar
- External Feedback
 - Anonymous reviewers of OOPSLA 2012 submission
 - Carlos Varela
 - Travis Dessell



Thank you!



```
import audience.questions.*
```



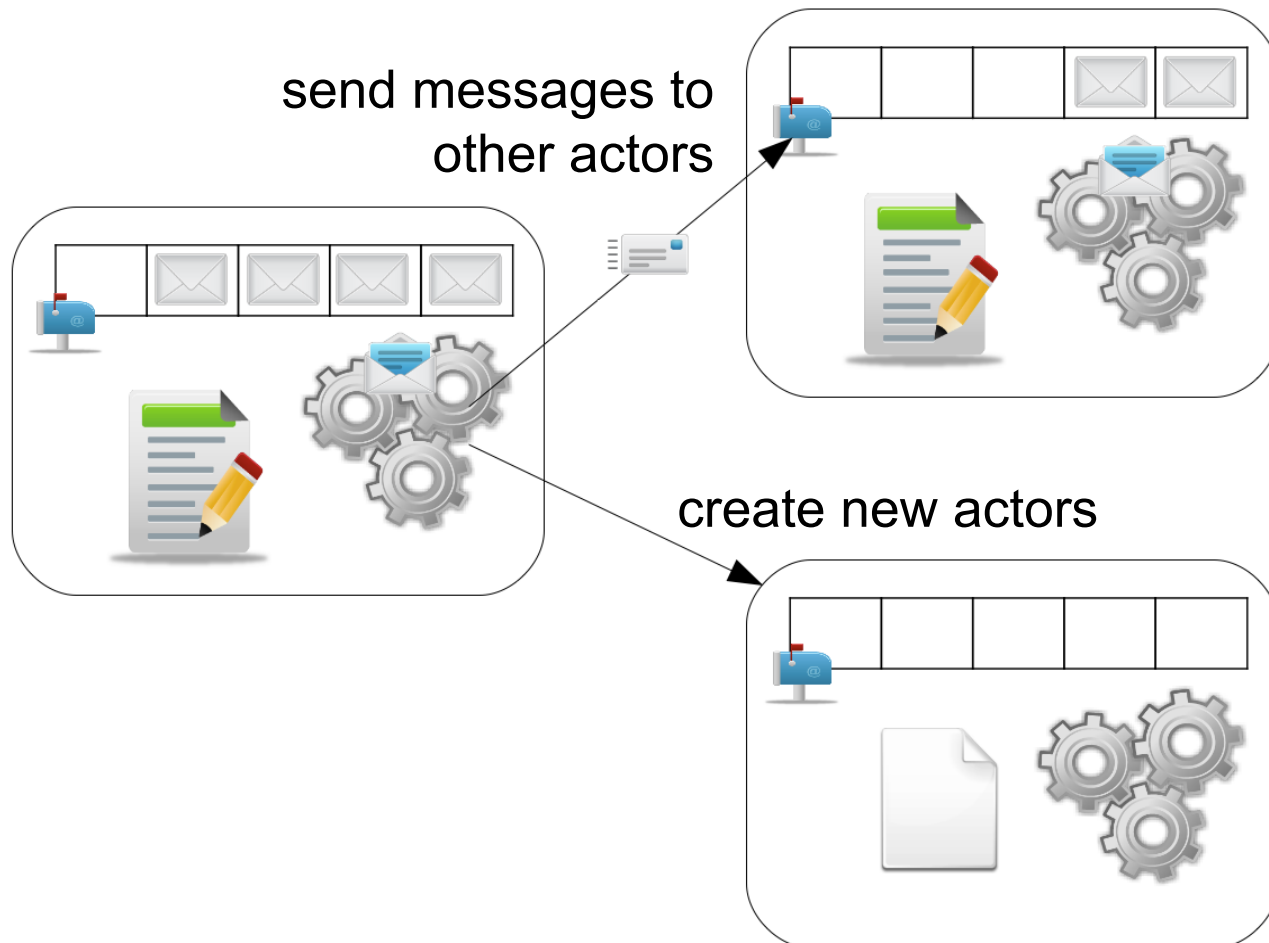


BACKUP SLIDES START HERE



Actor - Interactions

- Actors coordinate using *asynchronous* messaging
- Non-deterministic ordering of messages





Actor – Motivating Example



```
1  /** Scala code **/  
2  object FilterBankApp extends App {  
3    ...  
4    val sampler = ...  
5    val fir = new FirFilter(..., sampler).start()  
6    ...  
7    latch.await()  
8  }  
9  class FirFilter(..., nextStage: Actor)  
10     extends Actor {  
11    ...  
12    def act() = {  
13      loop { react {  
14        case FirItemMessage(value, coeffs) =>  
15          ...  
16          // compute the sum  
17          var sum = 0.0  
18          0 until coeffs.length foreach { index =>  
19            sum += buffer(index) * coeffs(index)  
20          }  
21          ...  
22          nextStage.send(DataItemMessage(sum))  
23          case ... => ...  
24        } } } }
```



Actors mapped to AFM



- Asynchronous messaging handled
- One message processed at a time invariant preserved
- Additional constructs used
 - *lingering* tasks
 - data-driven controls [for mailbox, see paper for details]
- No extra constraints placed on the Actors