



# Habanero-Java Library: a Java 8 Framework for Multicore Programming

PPPJ 2014  
September 25, 2014

Shams Imam, Vivek Sarkar

shams@rice.edu, vsarkar@rice.edu

Rice University



# Introduction

- Multicore processors are now ubiquitous
  - server, desktop, and laptop hardware
  - smaller devices: smartphones and tablets
- Parallelism is the future of computing
- Introduce parallelism early into the Computer Science curriculum



# Motivation for HJlib

- Writing parallel programs is hard
- Programmers need higher level parallel programming
- Distinguish between parallelism and concurrency
- Language approach requires extensions
  - Special compiler support
- Library-based approach integrates easily with existing code
  - Users can use IDE and tools of choice
  - Java 8 provides an excellent foundation for rich library support for parallelism and concurrency



# Contributions of this paper

- Habanero-Java library (HJlib) - a pure Java 8 library implementation of a multi-faceted task-parallel programming model
- EventDrivenControl API which can be used to add new parallel constructs to HJlib
  - All existing HJlib synchronization constructs (e.g., futures, data-driven futures, phasers) are also built using this API
  - Automatic support for AEM and deadlock detection
- Abstract Execution Metrics (AEM) framework for HJlib
- Deadlock detection tool for HJlib




# Simple Example: Two-way Parallel Array Sum

- Basic idea:
  - Decompose problem into two tasks for partial sums
  - Combine results to obtain final answer
  - Parallel divide-and-conquer pattern

Task 0: Compute sum of  
lower half of array

Task 1: Compute sum of  
upper half of array






# Simple Example: Two-way Parallel Array Sum

```
1.  // Start of Task T0 (main program)
2.  sum1 = 0; sum2 = 0; // sum1 & sum2 are static fields
3.  finish() -> {
4.      async() -> {
5.          // Child task computes sum of lower half of array
6.          for (int i=0; i < X.length/2; i++) sum1 += X[i];
7.      }); // end async
8.      // Parent task computes sum of upper half of array
9.      for (int i=X.length/2; i < X.length; i++) sum2 += X[i];
10. }); // end finish
11. // Parent task waits for child task to complete (join)
12. return sum1 + sum2;
```



# More complex HJlib example: Parallel Spanning Tree

```
1. class V {
2.     V [] neighbors; // Input adjacency list
3.     V parent; // Output spanning tree
4.     . . .
5.     boolean tryLabeling(final V n) {
6.         isolated(this, () -> {
7.             if (parent == null) parent = n;
8.         });
9.         return parent == n;
10.    } // end tryLabeling
11.    void compute() {
12.        for (int i=0; i<neighbors.length; i++) {
13.            final V child = neighbors[i];
14.            if (child.tryLabeling(this))
15.                //escaping async
16.                async(() -> { child.compute(); });
17.        }
18.    } // end compute
19. } // end class V
20. root.parent = root; // Use self-cycle to identify root
21. finish(() -> { root.compute(); });
```



**NOTE: this parallel structure cannot easily be expressed using standard Java constructs**



# HJlib Parallel Constructs

- Structured fork-join style parallelism ([async/finish](#))
- Parallel loops ([forall/forasync](#))
- Weak Atomicity ([isolated](#))
- Task Dependencies ([DataDrivenFuture](#))
- Point-2-Point Synchronization ([Phaser](#))
- Parallel Reductions ([FinishAccumulator](#))
- Asynchronous Message Passing ([Actor](#))
- Summary: HJlib supports an explicit parallel programming model at a higher level of abstraction and a wider range of parallel constructs than standard Java



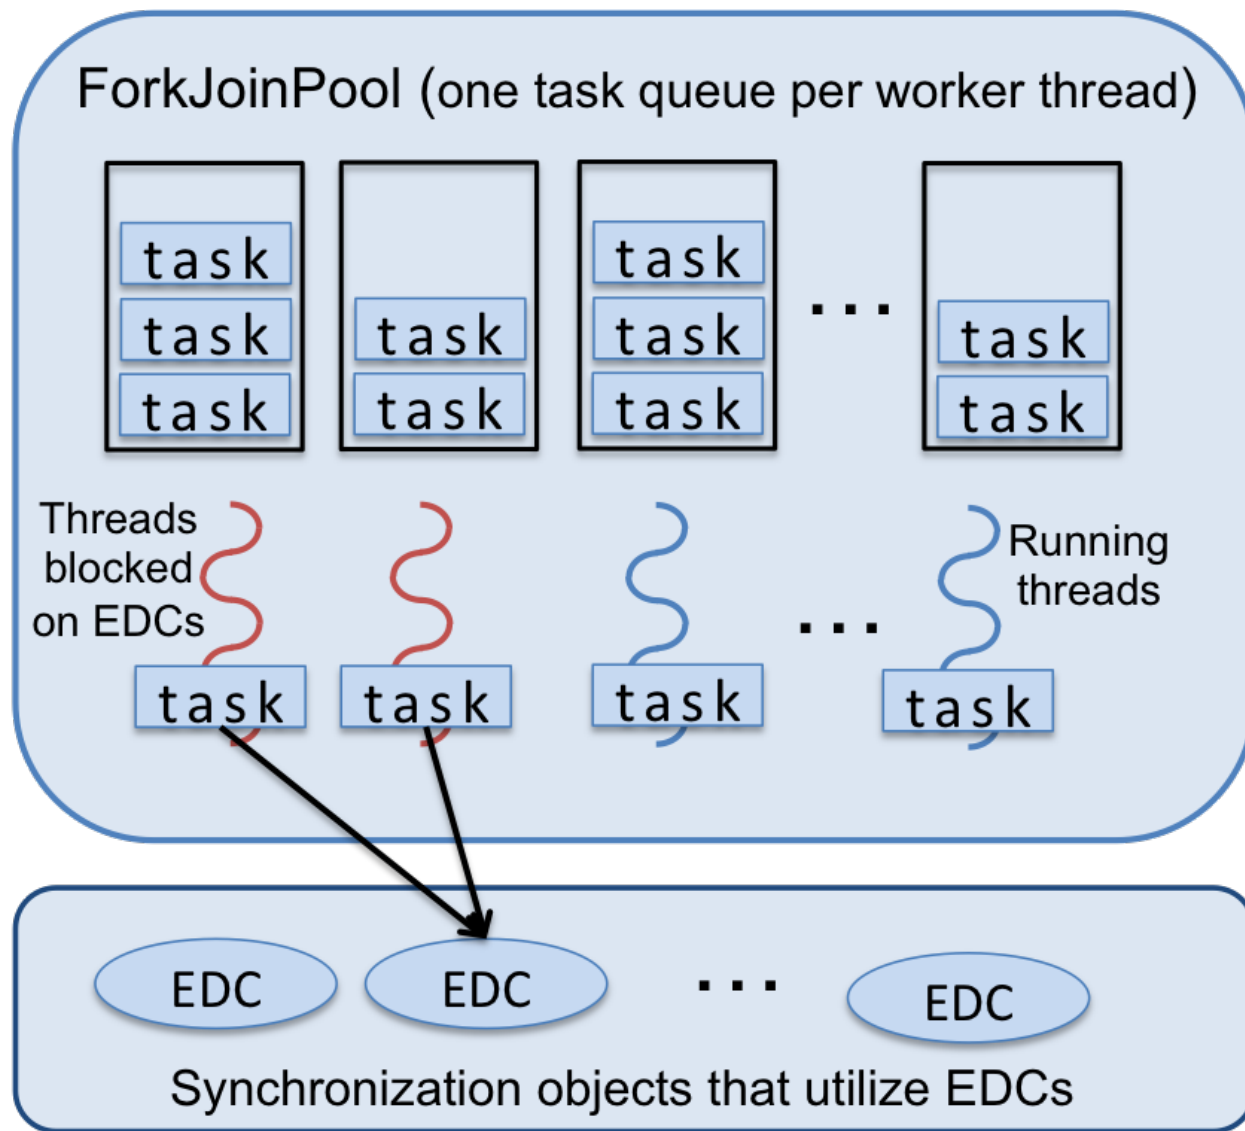


# Implementation

- No external dependencies
- Build on the capabilities offered by the Fork/Join Framework
  - ForkJoinPool work-stealing scheduler in JDK
- Built using Java 8 features
  - Lambda expressions
  - Functional interfaces
  - Older JVMs can be targeted by bytecode transformations
- Publisher/Subscriber model in runtime to add additional features



# Runtime Implementation






# Default Runtime

- Default runtime uses work-first policy with ForkJoin framework
- Blocking conditions block worker threads
  - Scheduler spawns additional worker threads to compensate
  - Runtime executes 'non-blocking' tasks before blocking
- All synchronizations implemented using EventDrivenControl data structure
  - Tracks which threads are blocked or can be resumed
- **NEW**: Alternative cooperative runtime is also available for HJlib
  - Pro: worker threads never block
  - Con: current continuation support adds overhead due to use of exceptions and on-the-fly bytecode transformation in class loader



# Event-Driven Control (EDC)

- Binds a value and a list of `java.lang.Runnable` blocks
  - Runnable blocks are code executed as callbacks
- Dynamic single-assignment of value (event)




The EDC is initially empty



# Event-Driven Control (EDC)

- Binds a value and a list of Runnable blocks
- Dynamic single-assignment of value (event)




Runnable blocks attach to the EDC and are not triggered until value is available (i.e. until event is satisfied)



# Event-Driven Control (EDC)

- Binds a value and a list of Runnable blocks
- Dynamic single-assignment of value (event)




Eventually, a value becomes available in the EDC  
(follows from deadlock freedom property of finish,  
futures, clocks, atomic)



# Event-Driven Control (EDC)

- Binds a value and a list of Runnable blocks
- Dynamic single-assignment of value (event)




This enables execution of runnable blocks attached to the EDC



# Event-Driven Control (EDC)

- Binds a value and a list of Runnable blocks
- Dynamic single-assignment of value (event)




Once value is available, subsequent  
runnable block attachment requests...





# Event-Driven Control (EDC)

- Binds a value and a list of Runnable blocks
- Dynamic single-assignment of value (event)



Synchronously execute the block  
(e.g. schedule a task into the work queue)



# Event-Driven Control API

- `currentTaskId()`:
  - returns a unique id of the currently executing task
- `newEDC()`:
  - factory method to create EDC instance
- `suspend( anEdc )`:
  - the current task is suspended if the EDC has not been resolved
  - Implementation attaches runnable block to resume task
- `anEdc.getValue()`
  - retrieves the value associated with the EDC
  - safe to call this method if execution proceeds past a call to `suspend()`
- `anEdc.setValue( aValue )`
  - resolves the EDC
  - triggers the execution of any EBs



# Parallel Constructs in Runtime

- Any task-parallel Synchronization Constraint can be supported.
  - Both deterministic and non-deterministic constructs
- All HJlib parallel constructs implemented using EDCs and event listeners (publish/subscriber) attached to runtime
- Key idea is to:
  - Translate the coordination constraints into producer-consumer constraints on EDCs
  - Block/Suspend consumers when waiting on item(s) from producer(s)
- Developers can add their implement own parallel constructs and add to HJlib
  - E.g. EventCount, others noted in future work




# Abstract Execution Metrics

- Enable users to reason about the performance of their parallel algorithms
- Compute total work done and critical path length
  - Dynamically generates the computation graph
  - Details for each construct in paper
  - Users can integrate custom parallel constructs
- Performance metrics are reproducible
  - Independent of physical machine used
  - Useful for debugging performance problems and comparing alternate implementations
- Computation graph can also be displayed visually



# Abstract Execution Metrics

- Prints the total Work and Critical Path Length (CPL)
- Supported for all HJlib constructs
- Enabled using:  
`HjSystemProperty.abstractMetrics.set(true)`
- Dump obtained by:  
`HjMetrics actualMetrics = abstractMetrics();`  
`AbstractMetricsManager.dumpMetrics(actualMetrics);`
- Can also use WORK and CPL metrics to obtain abstract time plots as shown on right






# Visual Graphs

- Example: MergeSort
  - Use each comparison operation as work

```
1 final int mid = M + (N - M) / 2;
2 finish(() -> {
3     async(() -> mergesort(A, M, mid));
4     async(() -> mergesort(A, mid + 1, N));
5 });
6 merge(A, M, mid, N);
```



# MergeSort Computation Graph





# Deadlock detector


- Enable users to debug their programs while using the various synchronization constructs
  - Reports diagnostic error message
- If they venture beyond the deadlock-free subset of Hjlib
  - `async/finish/future/phaser` programs are deadlock-free
- DDFs/Actors/EventCounts/Other custom constructs can cause deadlocks
- Algorithm relies on tracking the number of
  - ready tasks in the work queue
  - blocked or pending tasks (suspended on EDCs)
  - Deadlock when work queue is empty but there are pending tasks





# Deadlock Example

```
1 finish (() -> {  
2   HjDataDrivenFuture<Long> A = newDDF ();  
3   HjDataDrivenFuture<Long> B = newDDF ();  
4   asyncAwait(B, () -> A.put(B.get() + 3));  
5   asyncAwait(A, () -> B.put(A.get() + 5));  
6 });
```





# Pedagogic programming model

- Attractive tool for educators
- Educational resources available from COMP 322 website
  - Lecture notes and videos
- Extensive documentation and examples available
- Institutions can introduce parallel programming earlier in curricula
  - Based purely in Java
  - Garnered positive feedback from COMP 322 students
- Plan to use in MOOC version of COMP 322



# DEMO

- Setting up a simple HJlib project
- Some examples
  - ArraySum
  - QuickSort (Abstract Metrics)
  - Deadlock Detection
    - Data-Driven Futures
    - EventCount



# Future work, Ongoing Research



- Performant Cooperative runtime
- Visual Computation Graphs for AEMs
- Visual display of Deadlocks
- Parallel constructs
  - Selectors – an extension to actors
  - Eureka-style computations
- Use of HJlib for multicore parallelism in Android applications
- . . .
- See <http://habanero.rice.edu> for related research & publications



# Summary

- Pure library implementation on Java 8
- Introduces orthogonal parallel constructs with important safety properties
  - Simplifies parallel programming
- Feedback capabilities help the programmer debug applications
- Pedagogic: Attractive tool for both educators and researchers
- Educational material already available (COMP 322 lectures, videos, etc.)



# Questions

- Pure library implementation on Java 8
- Introduces orthogonal parallel constructs with important safety properties

```
import pppj.audience.Questions;  
import pppj.audience.Comments;  
import pppj.audience.Feedback;
```

- Simplifies parallel programming
- Feedback capabilities help the programmer debug applications
- Pedagogic: Attractive tool for both educators and researchers
- Educational material already available (COMP 322 lectures, videos, etc.)



# Backup-Slides



# Acknowledgments

- Rest of the Habanero Group
  - Vincent Cave
  - Jun Shirako
  - Sagnak Tasirlar
  - Jisheng Zhao








# Habanero-Java library

- Inspired from pedagogic Habanero-Java (HJ) language
- Emphasis on the usability and safety of parallel constructs
- Used in second-year undergraduate course at Rice

COMP 322: Fundamentals of Parallel Programming (Spring 2014)

Instructor:	<a href="#">Prof. Vivek Sarkar</a> , DH 3080	Graduate TA:	
	Please send all emails to <a href="#">comp322-staff</a> at <a href="#">rice dot edu</a>	Graduate TA:	
Assistant:	<a href="#">Benny Anderson</a> , <a href="#">anderson@rice.edu</a> , DH 3080	Graduate TA:	

- Actively used in multiple research projects at Rice



# Habanero-Java library

- Supports an explicit parallel programming model
  - A high level of abstraction
  - A wider range of parallel programming constructs
- A powerful and portable task parallel programming model
  - For the Java ecosystem
- Parallelize both regular and irregular applications



# Example: EventCount

- Keeps a count of the number of events in a particular class
- advance: signal the occurrence of an event
- await(v): suspends task until the value of the eventcount is at least v
- read: return count of the advance operations so far



# Example: EventCount

```
1 public final class EventCount {
2
3     Map<Long, EventDrivenControl> eventMap = new ←
        ConcurrentHashMap<>();
4     AtomicLong eventCounter = new AtomicLong(0);
5
6     public EventCount() {
7         EventDrivenControl edc = newEDC();
8         eventMap.put(0L, edc);
9         edc.setValue(Boolean.TRUE);
10    }
11    public void advance() {
12        long v = eventCounter.incrementAndGet();
13        eventMap.putIfAbsent(v, newEDC());
14        EventDrivenControl edc = eventMap.get(v);
15        edc.setValue(Boolean.TRUE);
16    }
17    public void await(final long v) {
18        eventMap.putIfAbsent(v, newEDC());
19        EventDrivenControl edc = eventMap.get(v);
20        suspend(edc);
21    }
22    public long read() {
23        return eventCounter.get();
24    }
25 }
```