

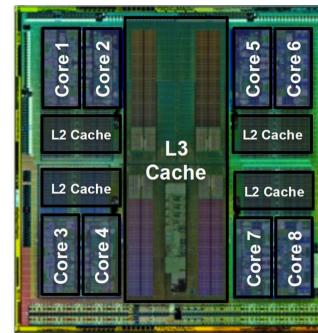
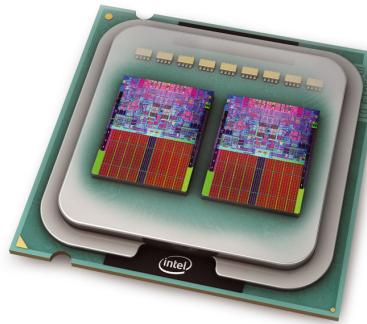
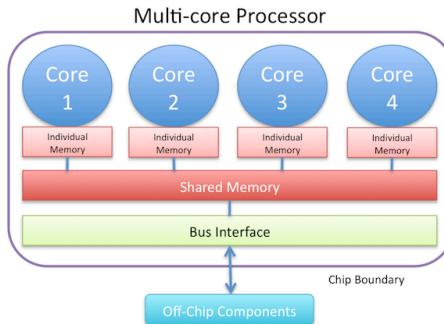
**Cooperative Execution
of
Parallel Tasks
with
Synchronization Constraints**

Shams Imam

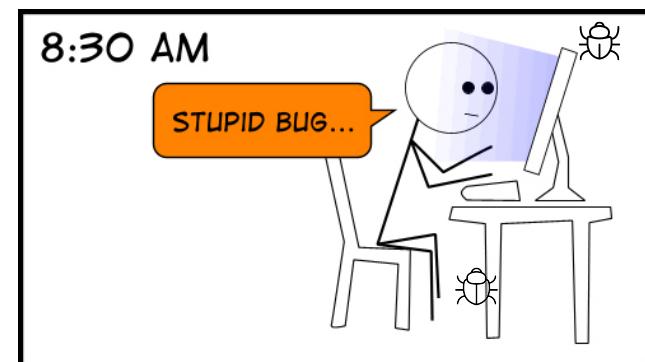
PhD Thesis Defense
Rice University
May 14, 2015

Introduction

- Multicore processors ubiquitous

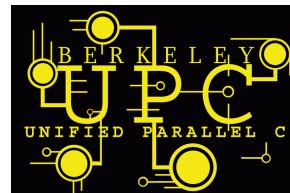


- Application performance requires parallelism
- Notoriously difficult to write performant parallel programs



Introduction...

- Parallel Programming models aid productivity
 - Offer Synchronization and Coordination constructs
 - E.g. Task Parallel Model, Actor Model



jetlang
Message based
concurrency for Java

ERLANG

IBM X10

OpenMP

 Microsoft
Asynchronous
Agents Library

Issue

- Current implementations for Synchronization and Coordination constructs are either
 - Not performant
 - Poor implementations **add overhead** in the runtime
 - O/S kernels designed to coordinate resource sharing inter-jobs rather than intra-jobs
 - e.g. future/barrier in Task Parallel Models
 - Not productive
 - **Complicate** user-level code by restricting synchronization patterns
 - e.g. barriers disallowed/behavior undefined inside OpenMP parallel loops, join only permitted on child tasks in Cilk, ...



Thesis Statement

“Many modern synchronization and coordination constructs in parallel programs can incur significant performance overheads on current runtime systems, or significant productivity overheads when the programmer is forced to restrict the synchronization patterns that they can use.

Our thesis is that that the use of cooperative scheduling techniques can help address the performance and productivity challenges of using modern synchronization and coordination constructs.

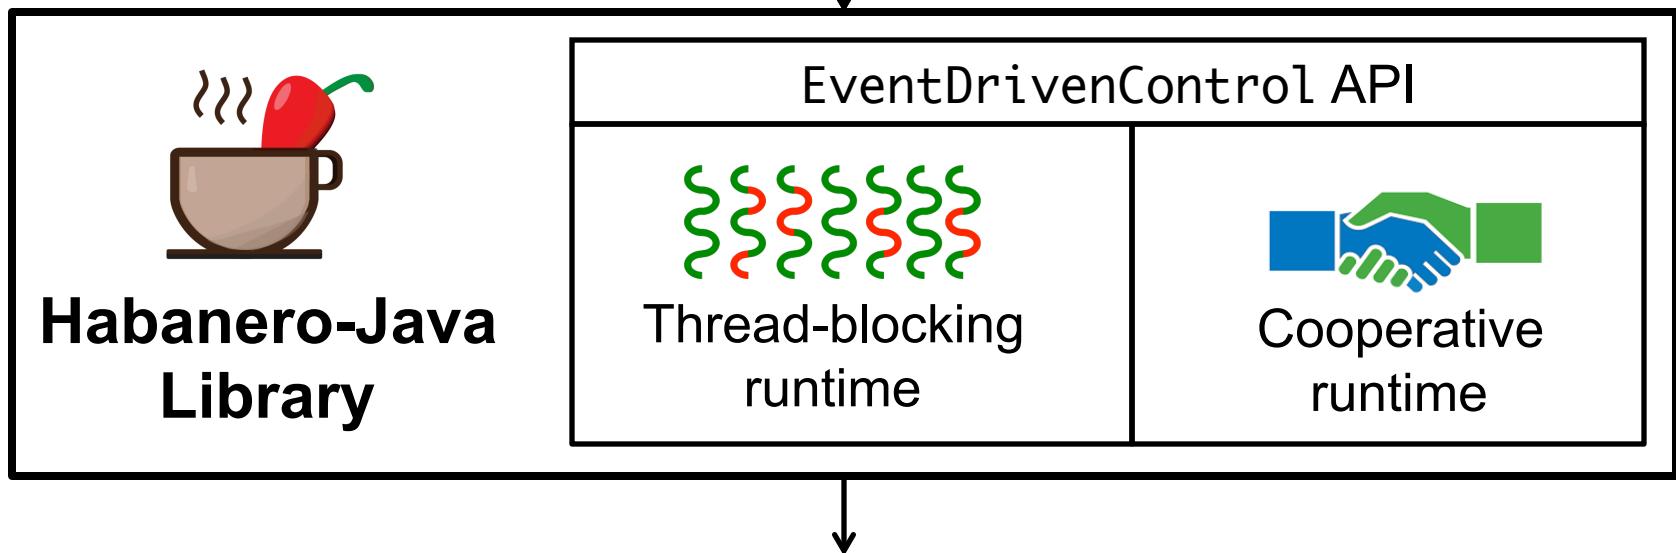
These techniques can be implemented using extensions to programming models, compilers, and runtime systems, depending on the desired trade-offs among performance, productivity, and support for legacy code.”

Thesis Overview

Programming Models



Runtimes

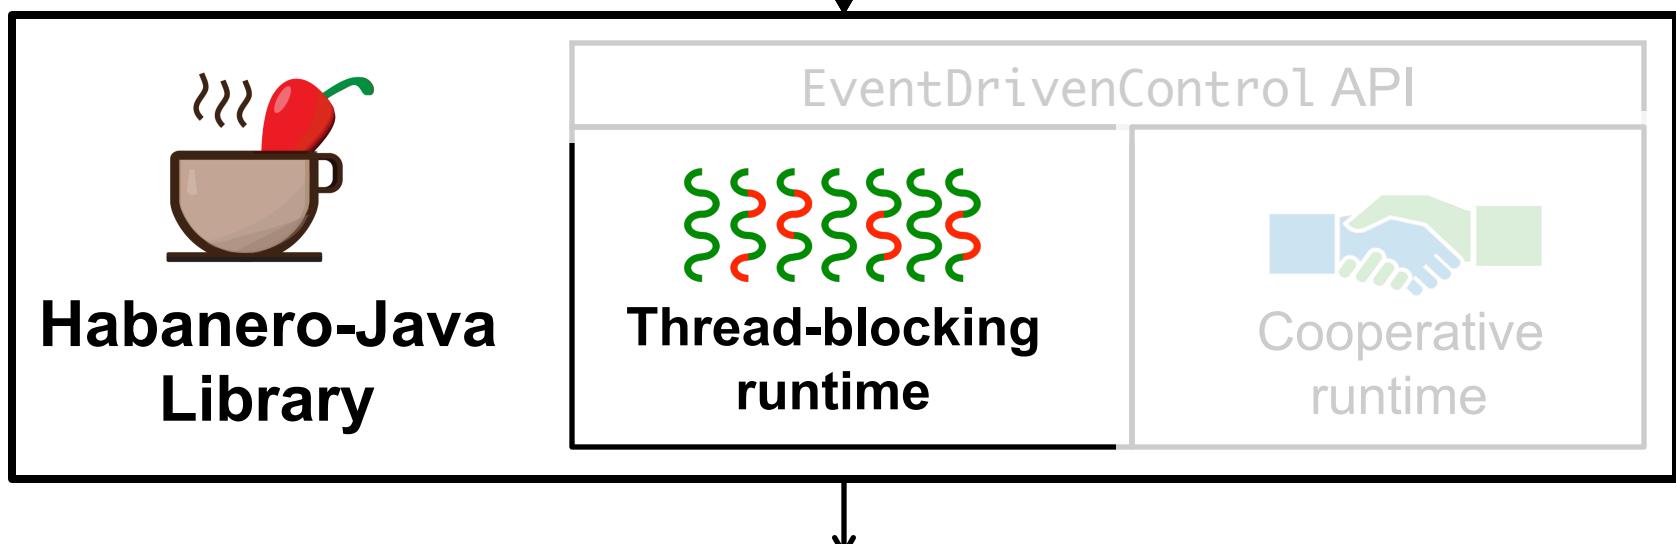


Habanero-Java Library

Programming Models



Runtimes





Habanero-Java Library (HJlib)

- Implementation of Habanero parallel programming model
- Library-based approach integrates easily with existing code
 - No special compiler support required
 - Java 8 lambdas simplify syntax
 - Thread-blocking runtime
 - Relies on JDK's work-stealing ForkJoinPool
 - Users can use IDE and debugger tools of choice
- Feedback capabilities
 - Helps programmer debug applications
 - Abstract metrics report
 - Deadlock detector

HJlib Contributions

- Higher level of abstraction with a wider range of parallel constructs than standard Java
- Pure **from-scratch** library implementation on Java 8
 - Built ground up with no external dependences
- Provides parallel programming framework for everyone
 - Implementation vehicle for concepts presented in this thesis
 - In use in multiple research projects
 - Pedagogic tool for educators

Habanero-Java Library: a Java 8 Framework for Multicore Programming. **Shams Imam**, Vivek Sarkar. 11th International Conference on the Principles and Practice of Programming on the Java Platform (**PPPJ'14**), September 2014.

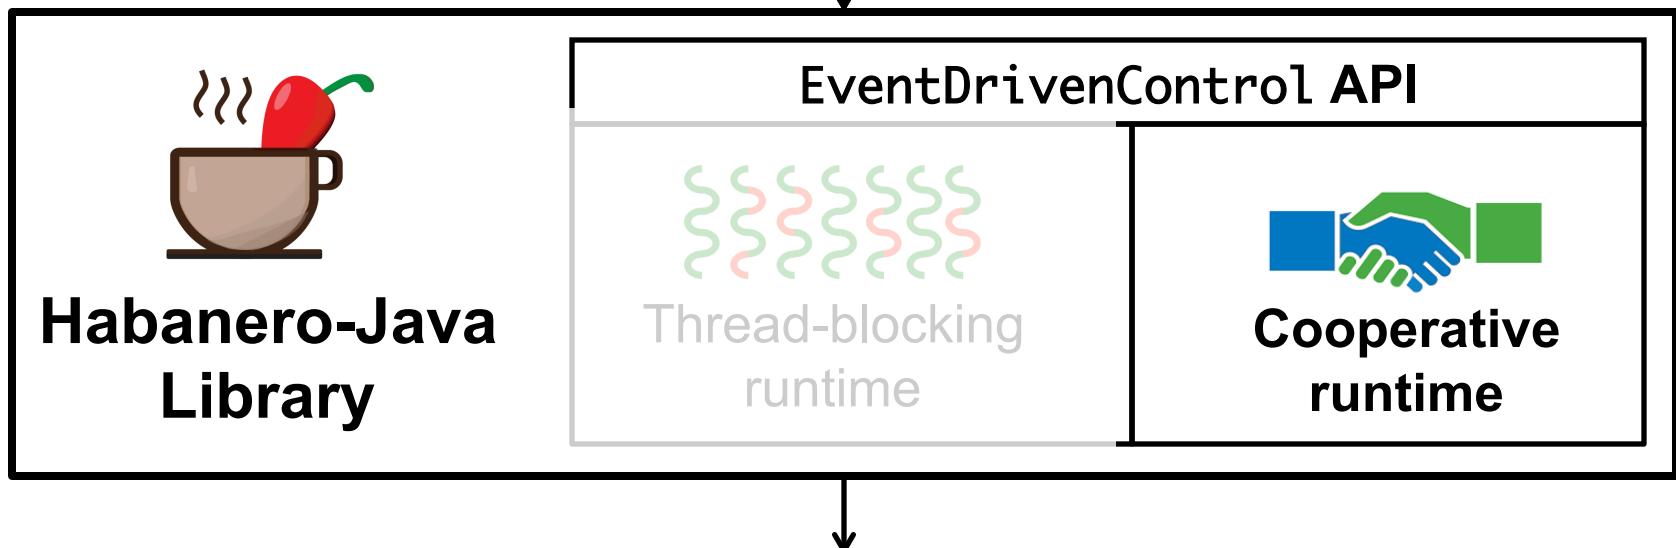


Cooperative Runtime for Parallel Tasks

Programming Models



Runtimes



Task-Parallel Model (TPM)



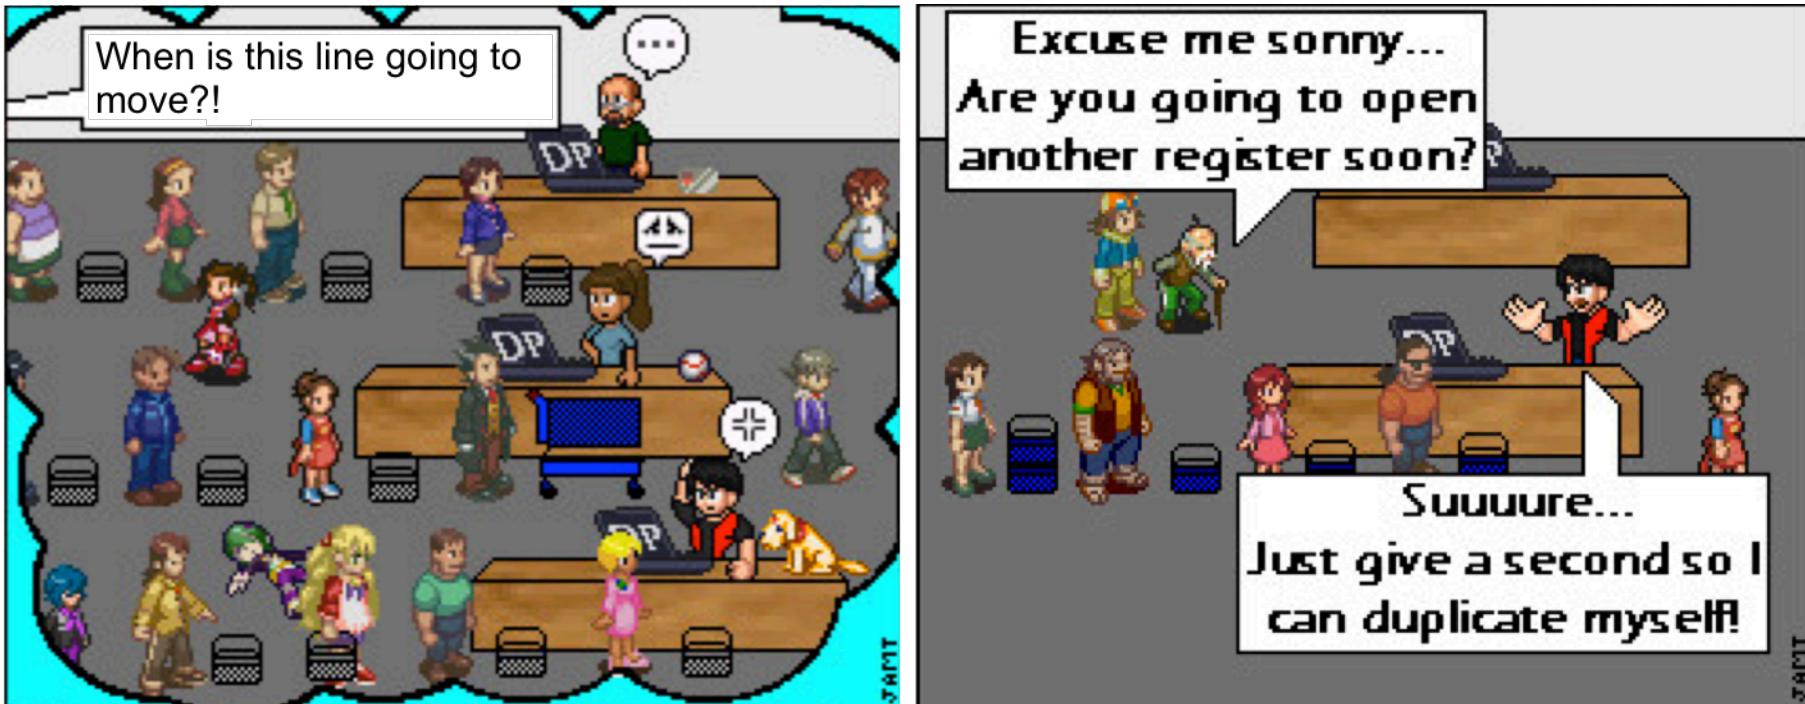
- Tasks, Work Queues, and Worker Threads
- Runtime manages load balancing and synchronization

Synchronization Constraints

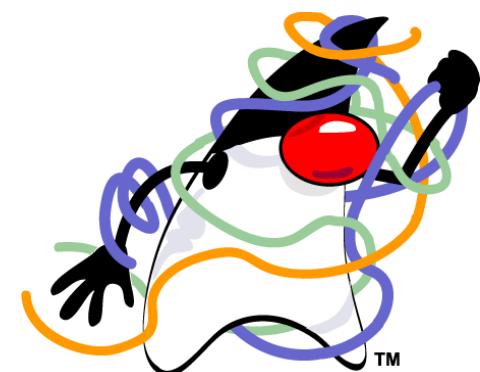


- Dependences between tasks
- Prevent an executing task from making further progress
 - Needs to synchronize with other executing task(s)

Common TPM Runtime Solution: Spawn Additional Worker Threads

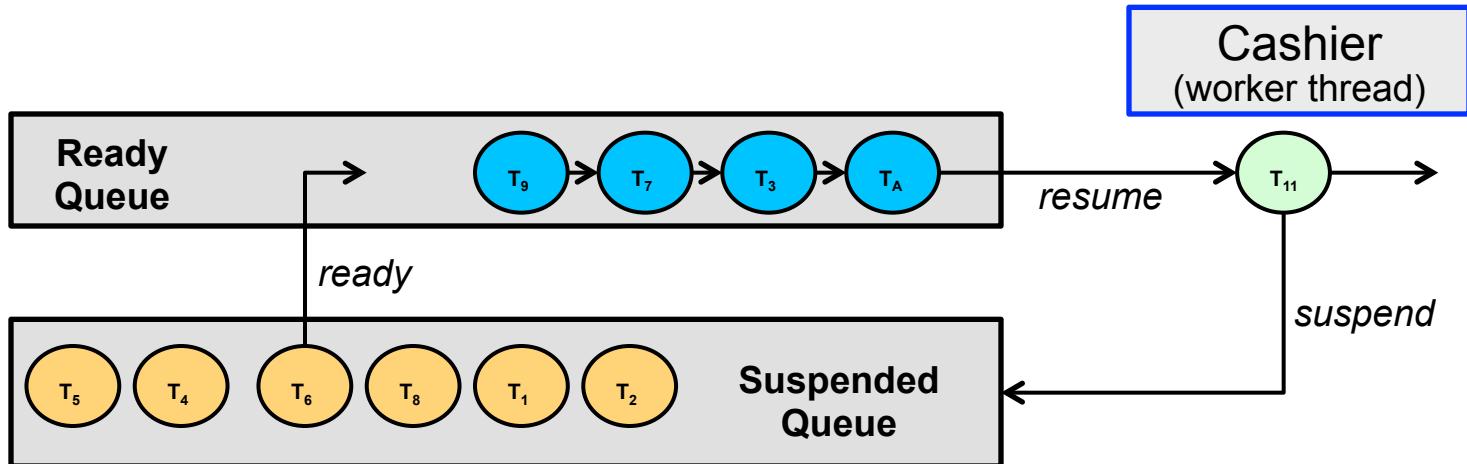


- Each thread needs its own system resources
- Exhaustion of memory or other system resources
- Scheduler tangled in thread scheduling decisions
- **Thread-blocking approaches do not scale!**



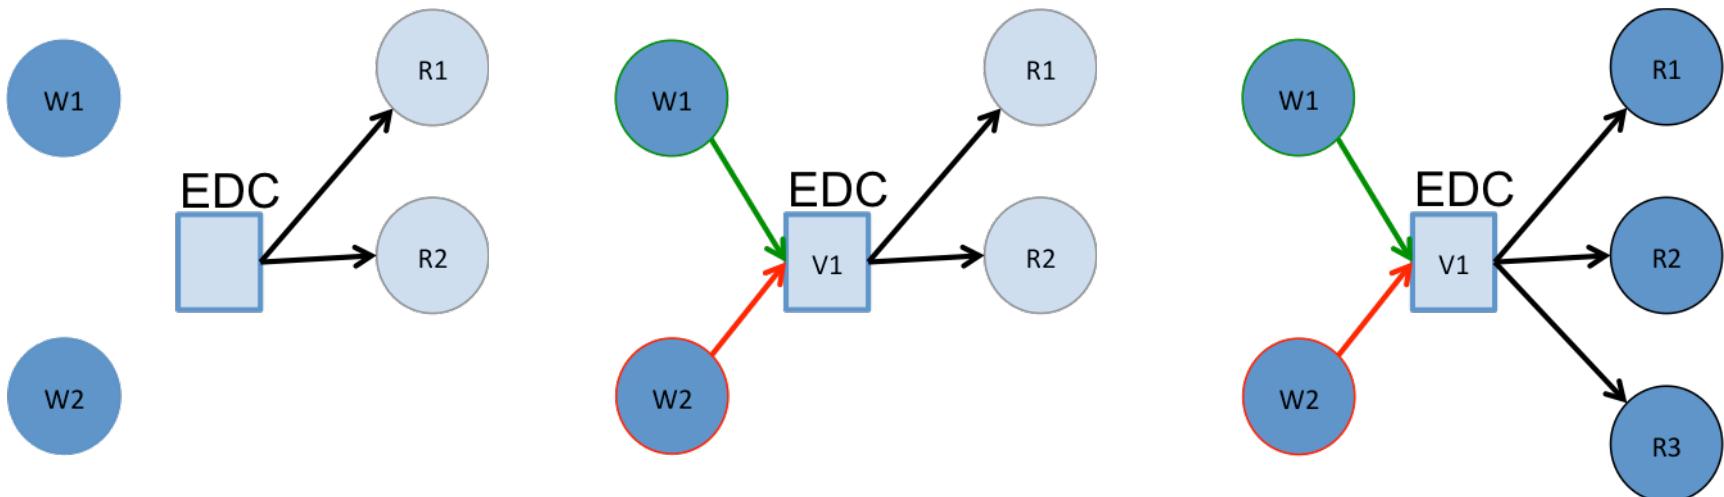
Proposed Solution: Cooperative Scheduling

- A **Cooperative** Approach is more efficient
- Task suspended by runtime at synchronization points
 - Task Continuation (TC) moved to the suspended queue
 - Compiler support needed to generate the continuations
- TC moved to ready queue when the task can make progress
- TC resumed when executed on worker thread



Event-Driven Control (EDC)

- Delays execution of readers until value available
- Dynamic single-assignment of value (event) by writer



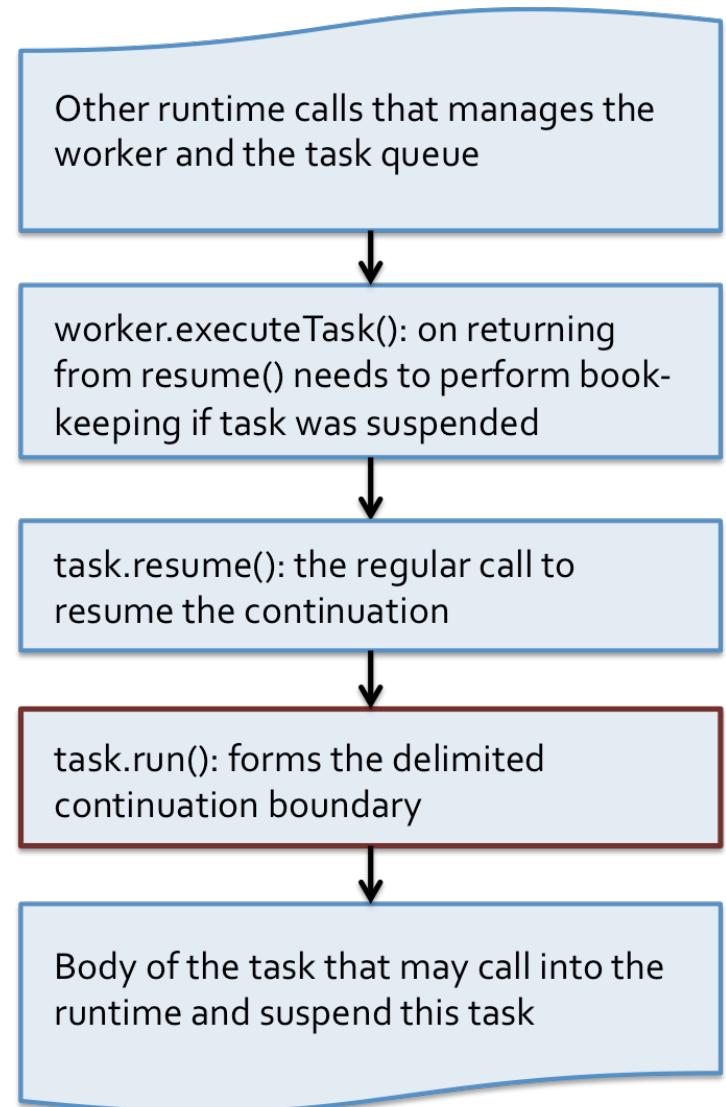
1. Initially empty,
readers suspend

2. Some writer succeeds,
value becomes available

3. All readers execute,
subsequent readers
do not get suspended

Cooperative Runtime – Call Stack

- **Help-first policy**
 - Task has a stack of its own
 - Task can be executed by any of the worker threads
- Task wrapped to form One-shot Delimited Continuations
- Worker thread manages when tasks get
 - **cooperatively** added to suspended queue
 - removed from ready queues for execution by worker threads



Cooperative Runtime

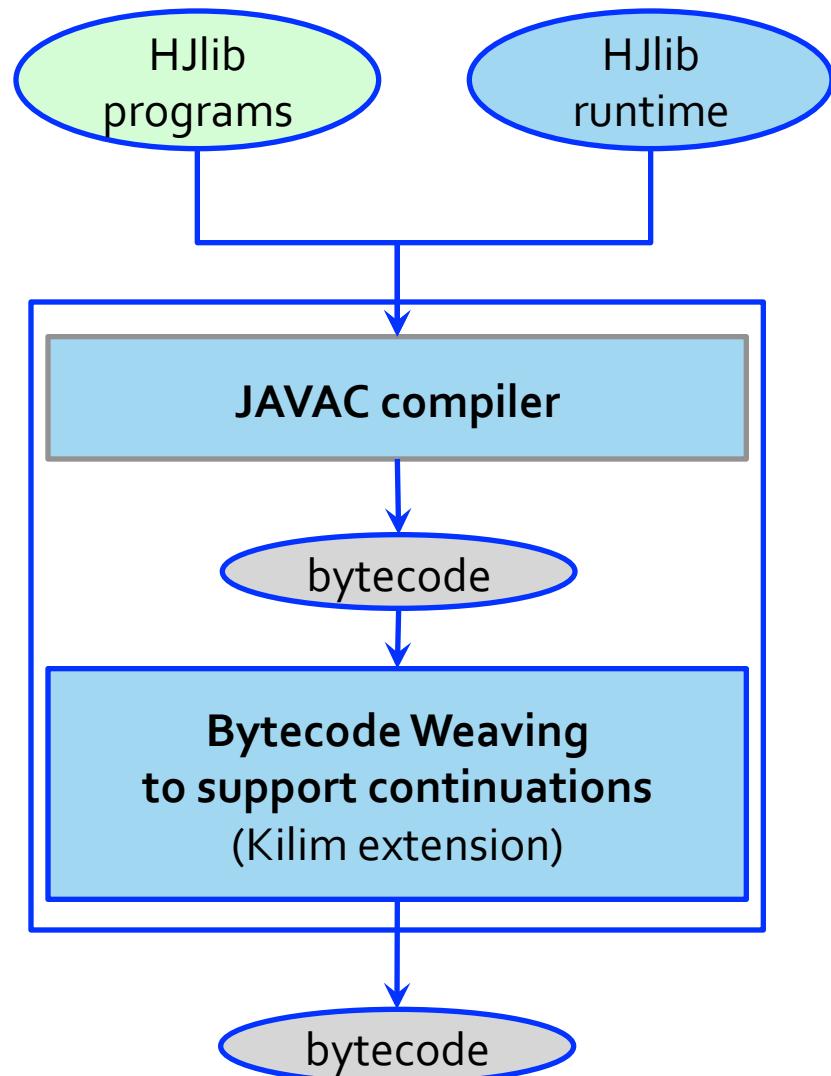
- We expose EDCs as an internal API in our runtime
 - Read / Write / Query on value
 - Suspend till value becomes available
- Continuations not exposed to developer
 - Notorious for being hard to use and to understand
- Developers write thread-based code
 - Compiler handles code transformations to support continuations
 - One-shot delimited continuations implemented more efficiently than general continuations

Reducing Overheads compared to Thread-Blocking

- Creating additional worker threads is contradictory to goal of TPM
- 1100 ns per thread context switch (without cache effects)
- In contrast
 - Object allocation takes around 30 ns
 - Method call takes around 5 ns
 - Setting fields takes around 1 ns
- Key assumption of this work
 - Creating threads and context switching is expensive
 - Task creation and allocating synchronization objects are cheaper

Implementation: Compiler Infrastructure

- Standard Java code as input
- First compilation phase uses javac compiler to generate bytecode
- Second Compilation phase CPS transforms code to support continuations
 - This phase can also be done during class-loading time
- Output is bytecode that runs on any standard JVM

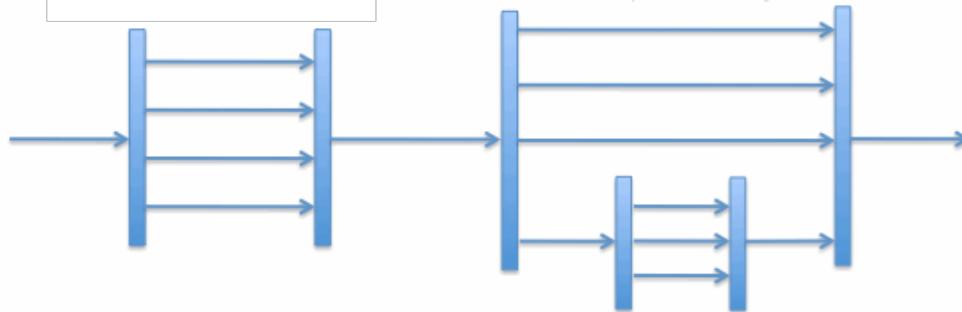


Synchronization Constructs

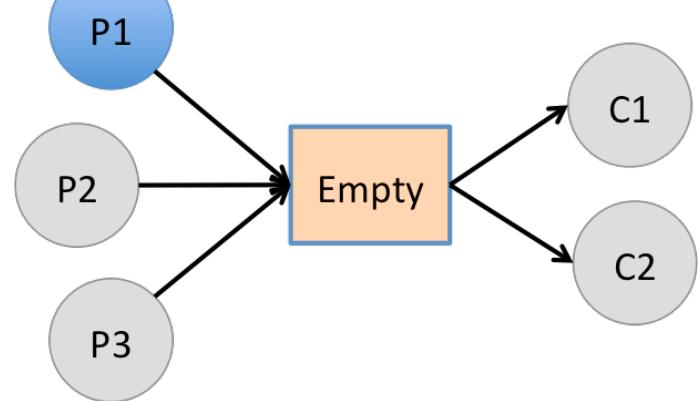
- Key idea is to:
 - Translate the synchronization constraints into producer-consumer constraints on EDCs
 - Use Delimited Continuations to suspend consumers when waiting on item(s) from producer(s)
- Any task-parallel Synchronization Constraint can be supported.
 - Both deterministic and non-deterministic constructs
 - Including `async-finish`, `futures`, `barriers`, `atomic/isolated` and `actors`

Synchronization constructs...

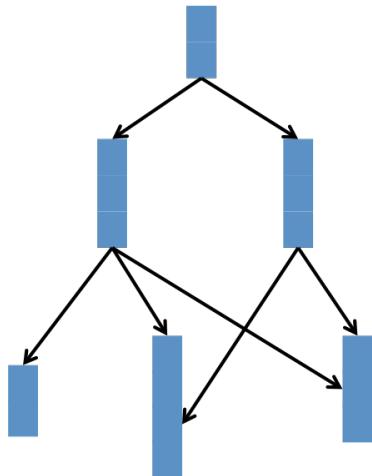
Nested fork-join



Synchronization Variable



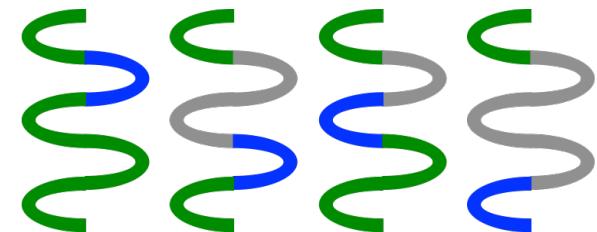
Futures



Barriers



Critical-Section blocks



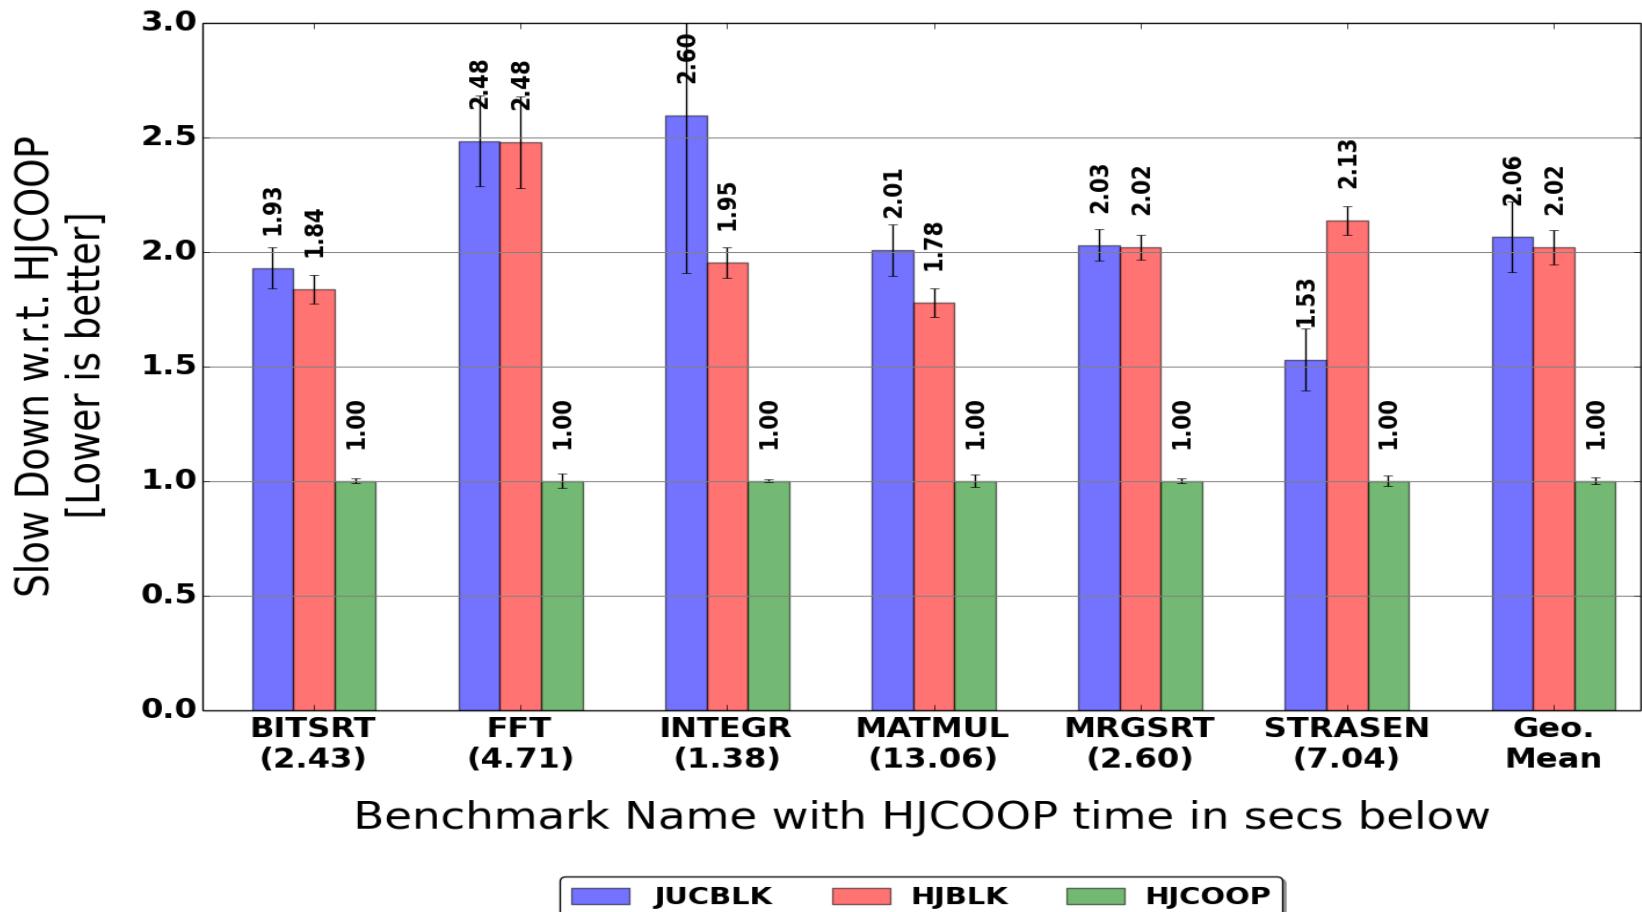
- Recipes for these and other constructs described in thesis...



Experimental Setup

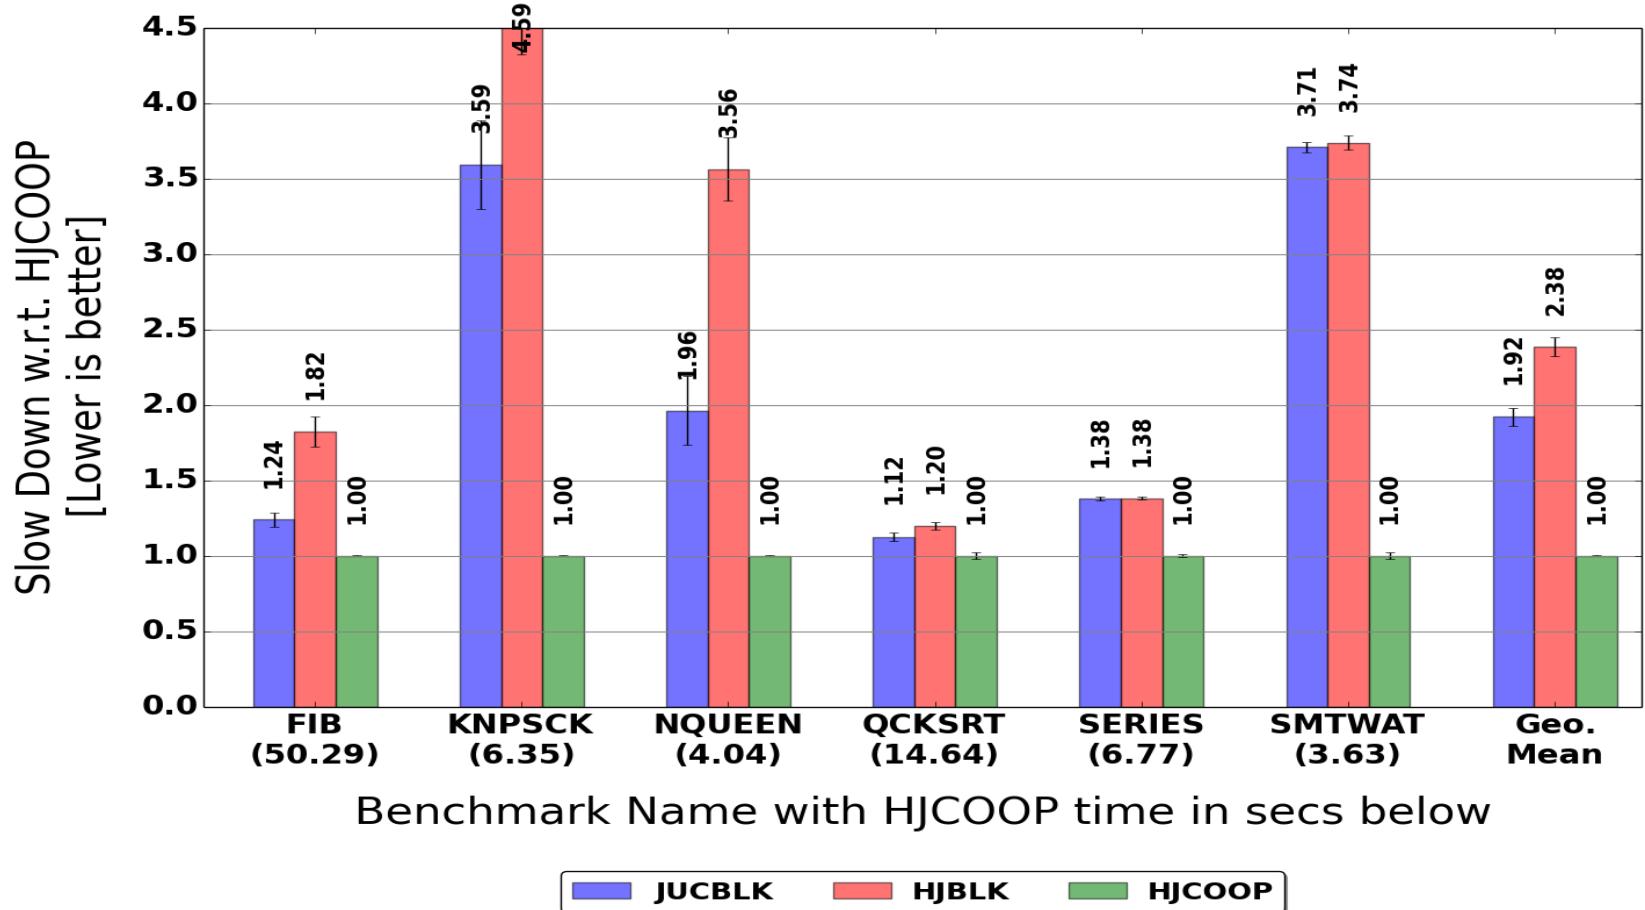
-  Four 8-core 3.8 GHz IBM POWER7
-  256 GB of RAM
-  32 KB L1 Cache
- Software
 -  IBM Java SDK Version 1.8.0
 -  Habanero-Java library 0.1.2
 - Thread-blocking scheduler
 - Cooperative scheduler enabled via option
- Benchmarks run with single place
 -  32 worker threads per place
 - 64 GB memory allocated to JVM
 -  Mean of best 30 out of 100 execution times reported

async-finish Benchmarks



- Benchmarks that create lots of dynamic finish scope instances chosen
- Thread-blocking runtimes crash on larger inputs

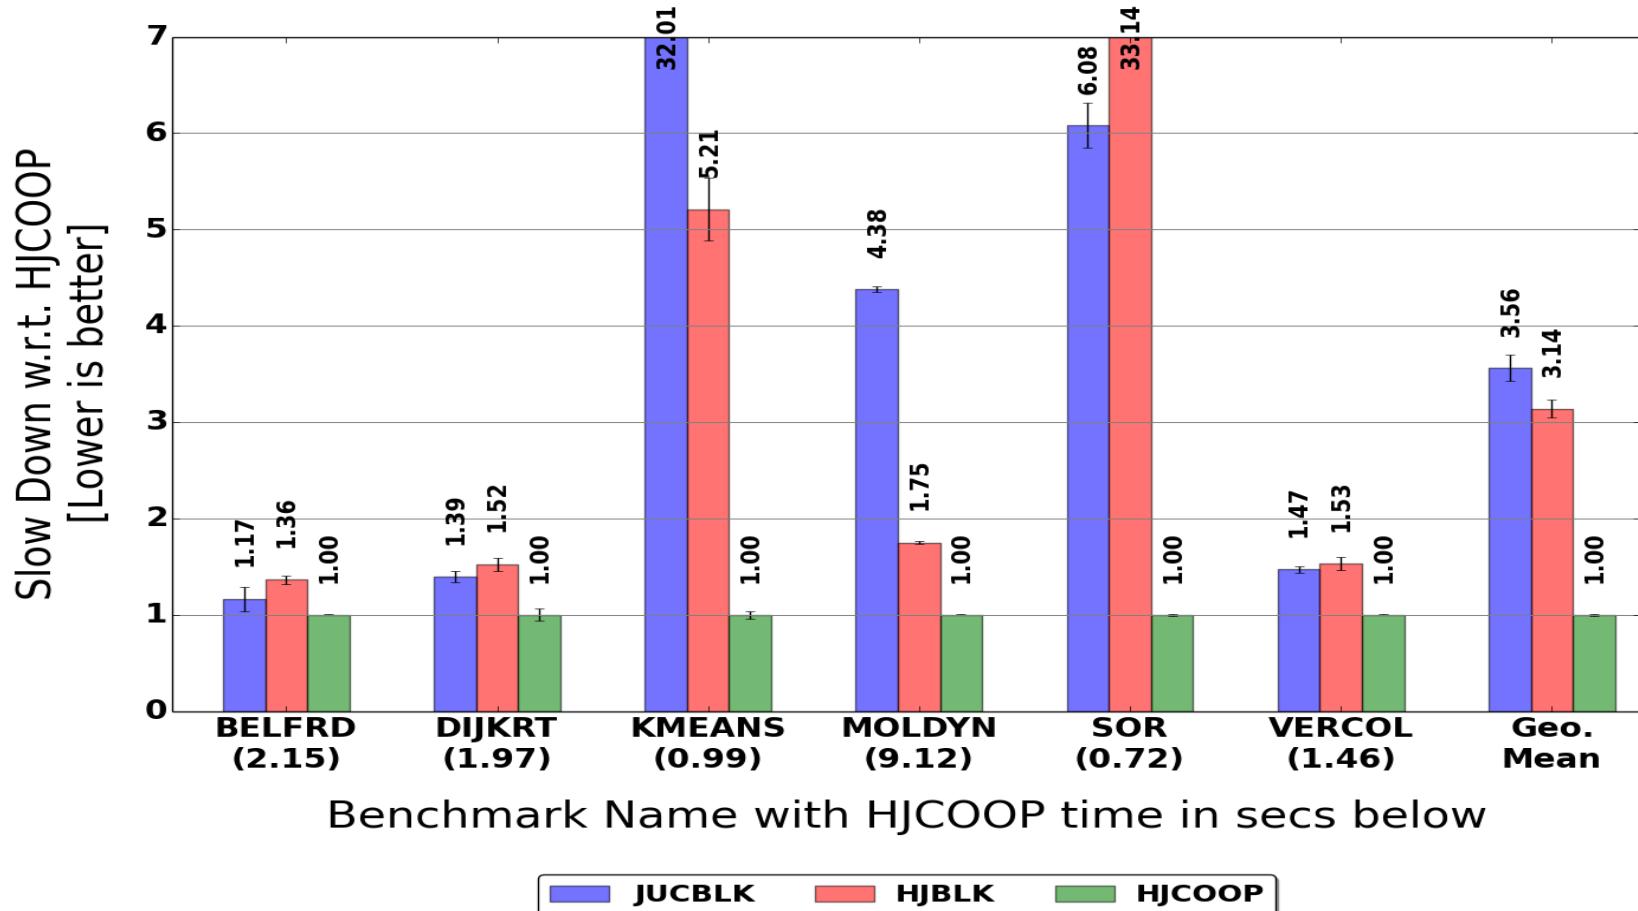
future Benchmarks



- Benchmarks that eagerly attempt to resolve futures chosen
- Thread-blocking runtimes crash on larger inputs



Phaser Benchmarks



- Benchmarks chosen such that more than 32 tasks are registered on phaser

Related work

- Qthreads [IPDPS '08]
 - Continuations with lightweight call stack stitching
 - Stack size configured manually
- Glasgow Haskell Compiler [Haskell '07]
 - Provides continuation support directly
 - Uses polling to resume continuations
- C++ implementation of X10 [PPoPP '12]
 - Work-first policy with dedicated implementation for each construct
 - Implementation supports async-finish and futures
 - Implementation did not support clocks (phasers)

Cooperative Runtime Contributions

- **Event-Driven Controls API**
 - Can support **any** task-parallel synchronization construct
 - Recipes for various task-parallel synchronization constructs
- **Cooperative** runtime for scheduling tasks
 - Transparent use of One-shot Delimited Continuations
 - Supports triggering the **enablement of multiple** suspended tasks
- **Performs better** than runtimes that use thread-blocking

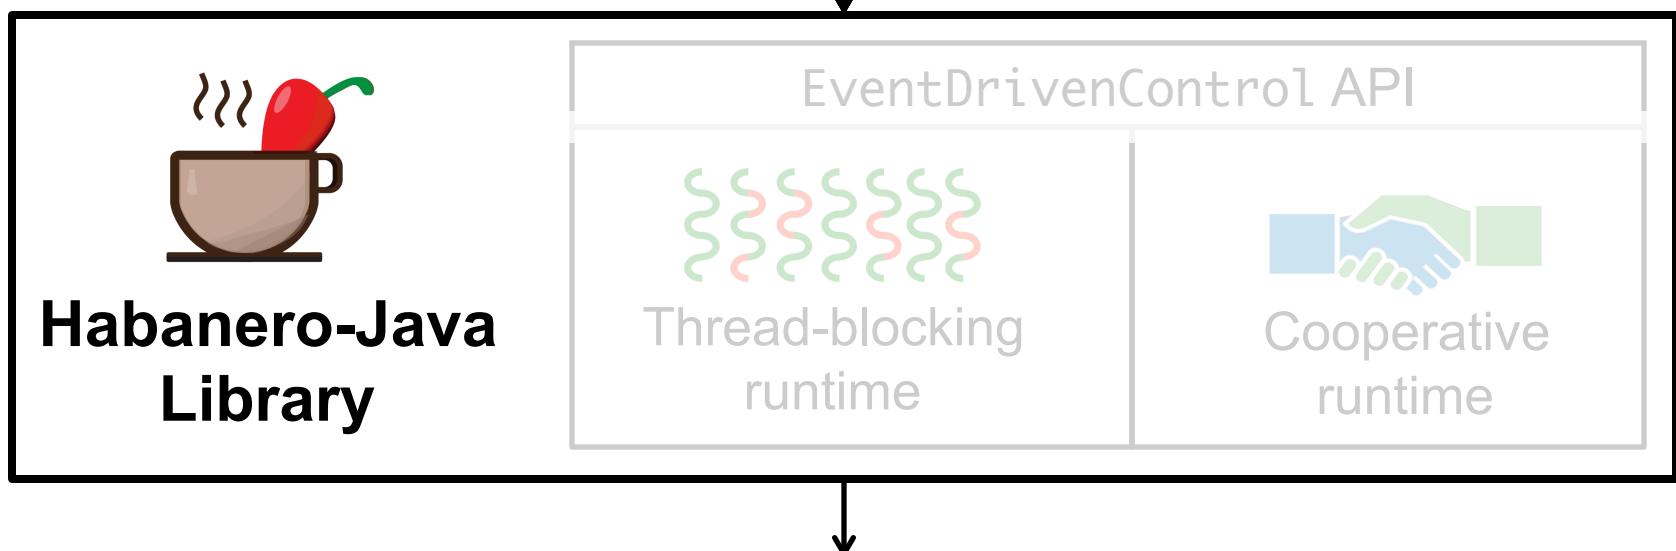
Cooperative Scheduling of Parallel Tasks with General Synchronization Patterns. **Shams Imam**, Vivek Sarkar. 28th European Conference on Object-Oriented Programming (**ECOOP**), July 2014.

Eureka Programming Model

Programming Models



Runtimes



Eureka Style Computation

- Mostly optimization and search problems
- Speculative parallelism abundant
 - Results may or may not be needed
 - Each task attempts to find a result
- Announces result has been found
 - An "aha!" moment – **Eureka** event
 - Can curtail rest of the computation



Simple Approach: Matrix Search

```
6  def doComputation(matrix, goal) {  
7      val token = atomicRefFactory()  
8      finish  
9      for rowIndices in matrix.chunks()  
10         async  
11             for (r in rowIndices)  
12                 processRow(matrix(r), r, goal, token)  
13             // return either [-1, -1] or a valid index [i, j]  
14         return token.get()  
15     }  
16     def processRow(rowData, r, goal, token) {  
17         for (c in rowData.indices())  
18             if goal.matches(rowData(c)) // eureka!!!  
19                 token.set([r, c])  
20             return  
21     }
```

Related Work:

Function-Scope termination

- Cancellation at lexical scope of task creation
- E.g. Cilk abort, OpenMP 4.0

```

1  async
2    for (r in rowIndices)
3      processRow(matrix(r), r, goal, token)
4      check(token) // check cancellation
      } } Can trigger cancellation here

6  def processRow(rowData, r, goal, token) {
7    for (c in rowData.indices())
8      if goal.matches(rowData(c))
9        token.set([r, c])
10       return
11   }
      } } Cannot trigger cancellation in this lexical block

```

- Cancellation not possible from nested function call!

Related Work: Exceptions for Control Flow

- Throw exceptions to terminate tasks
- Only edit specific program points
- E.g. MS Task Parallel Library, Java Threads

```

1  async
2      for (r in rowIndices)
3          try
4              processRow(matrix(r), r, goal, token)
5          catch (ex: Exception)
6              exceptionHandling(ex) ← EurekaException
7
8      def processRow(rowData, r, goal, token) {
9          for (c in rowData.indices())
10         if goal.matches(rowData(c))
11             token.set([r, c])
12             throw EurekaException()
13     }

```

- User may inadvertently catch exception

Related Work:

Manual termination via Cancellation Tokens

- Manual periodic checks with returns
- User controls responsiveness

```

1  async
2    for (r in rowIndices)
3      if token.eureka()
4        return
5      processRow(matrix(r), r, goal, token)

7  def processRow(rowData, r, goal, token) {
8    for (c in rowData.indices())
9      if token.eureka()
10     return
11     if goal.matches(rowData(c))
12       token.set([r, c])
13   }

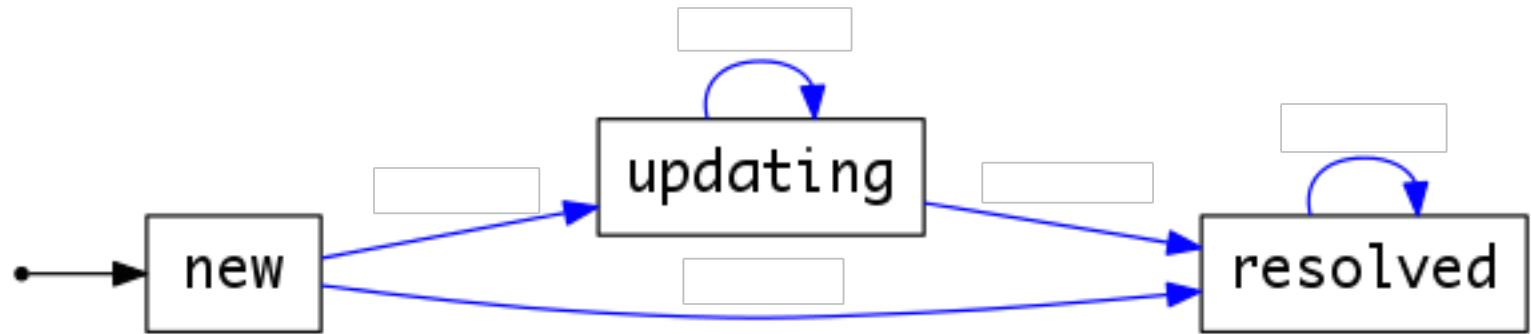
```

Repeated checks
which are written
manually

- Cumbersome to write
 - Pass-through methods need instrumentation
- Impossible to support inaccessible functions

Detecting eureka events

- **Our solution:** Eureka construct to abstract various patterns
- Simplifies grouping of speculative tasks



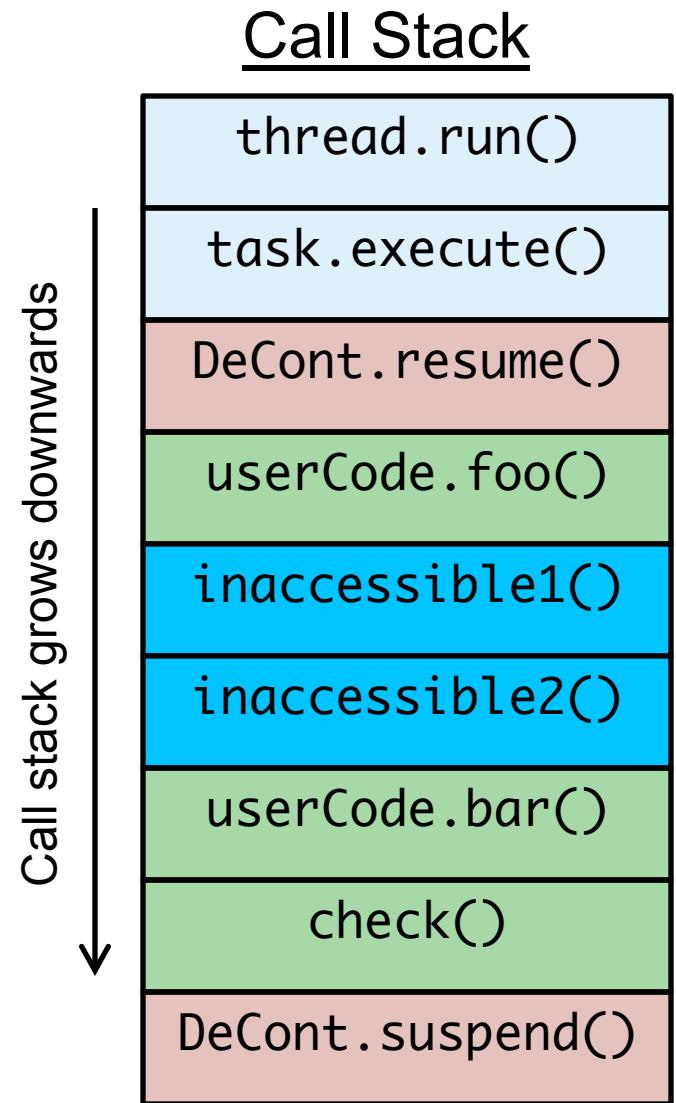
- NEW: freshly created
- UPDATING: changeable state
- RESOLVED: stable state
- Desired invariant: Monotonicity in state transitions
- Once resolved, eureka always stays resolved

Eureka Programming Model

- **finish (eureka) S1**
 - Multiple finishes can register on same Eureka
 - Wait for all tasks **or** until eureka is resolved
- **async**
 - Launches an asynchronous task
 - Inherit Eureka registration from immediately-enclosing finish
- **offer()**
 - Triggers eureka event on registered Eureka
- **check()**
 - Causes task to terminate if Eureka resolved

Task Termination Technique

- **Our solution:** Use Delimited Continuations!
 - Explained earlier
 - Even cheaper – no need to save state during suspension
- Rely on cooperative task termination techniques
 - **check** status of eureka construct
 - **terminate** task by suspending task



Eureka Version of Example

```
5  def doComputation(matrix, goal) {
6      val eu = eurekaFactory()
7      finish(eu) // eureka registration
8      for rowIndices in matrix.chunks()
9          async
10             for r in rowIndices
11                 processRow(matrix(r), r, goal)
12     return eu.get()
13 }
14 def processRow(rowData, r, goal) {
15     for c in rowData.indices()
16         check([r, c]) // cooperative termination check
17         if goal.matches(rowData(c))
18             offer([r, c]) // trigger eureka event
19 } }
```

Eureka Variants

- Both deterministic and nondeterministic computations
- Examples
 - Parallel Search
 - K-Count Eureka
 - N-Version Programming
 - Branch and Bound Optimization
 - Soft Real-Time Deadlines
 - Convergence Iterations
 - Binary Composition of variants
 - Nesting computations
 - ...

Eureka Variants...

```
def eurekaFactory() {
    val initialValue = [-1, -1]
    return new SearchEureka(initialValue)
}
```

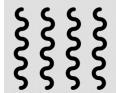
```
def eurekaFactory() {
    val K = 4
    return new CountEureka(K)
}
```

```
def eurekaFactory() {
    // comparator to compare indices
    val comparator = (a, b) -> {
        ((a.x - b.x) == 0) ? (a.y - b.y) : (a.x - b.x)
    }
    val initialValue = [INFINITY, INFINITY]
    return new MinimaEureka(initialValue, comparator)
}
```

```
def eurekaFactory() {
    val time = 4.seconds
    return new TimerEureka(time)
}
```

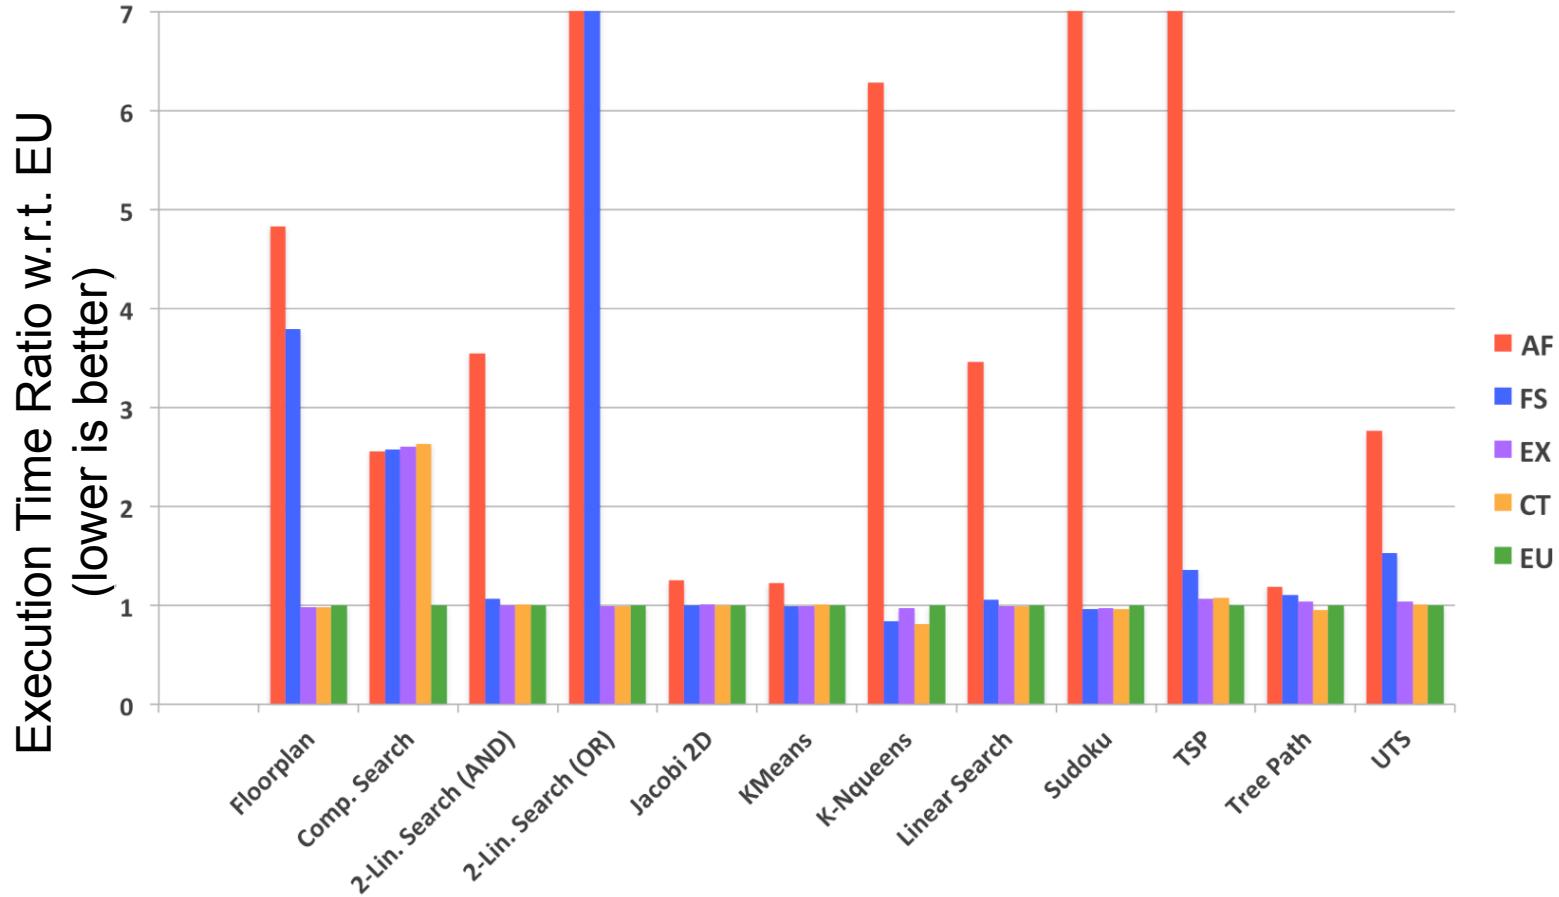
```
def eurekaFactory() {
    val units = 400
    return new EngineEureka(units)
}
```

Experimental Setup

-  Four 8-core 3.8 GHz IBM POWER7
-  256 GB of RAM
-  32 KB L1 Cache
- Software
 -  Habanero-Java library 0.1.4
 - Cooperative scheduler that uses continuations
 - Other variants implemented from scratch
- Benchmarks run with single place
 -  32 worker threads per place
 - 64 GB memory allocated to JVM
 -  Mean of best 50 out of 150 execution times reported

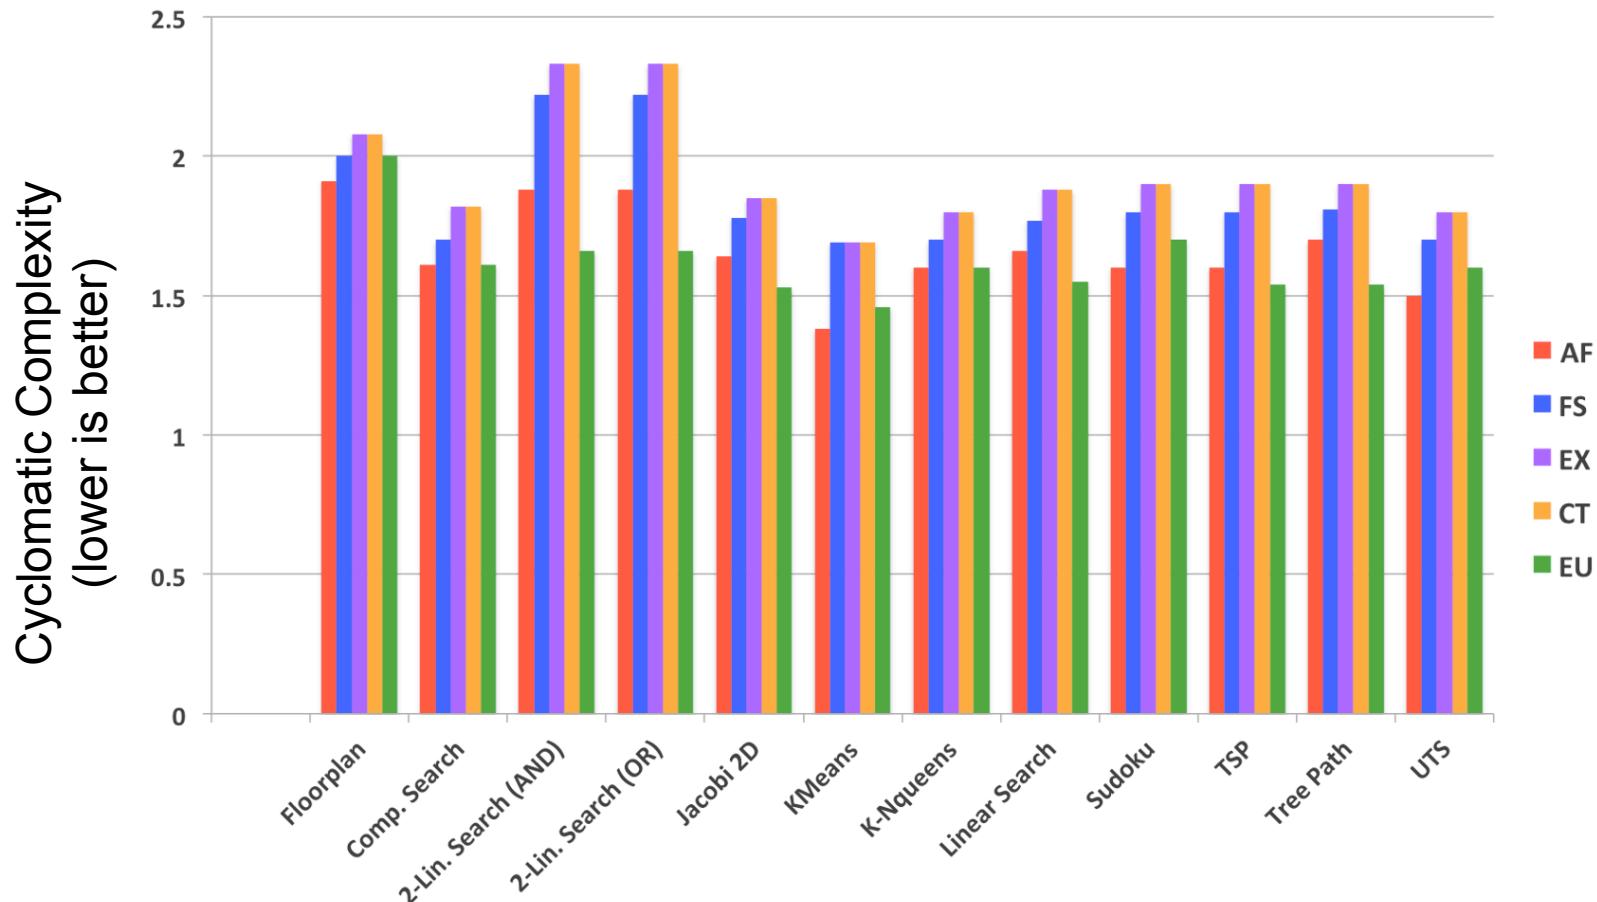
Performance Results

- Legend: AF: async-finish; CT: Cancellation-token; EX: Exception-based; FS: Function-scoped; EU: Eureka.



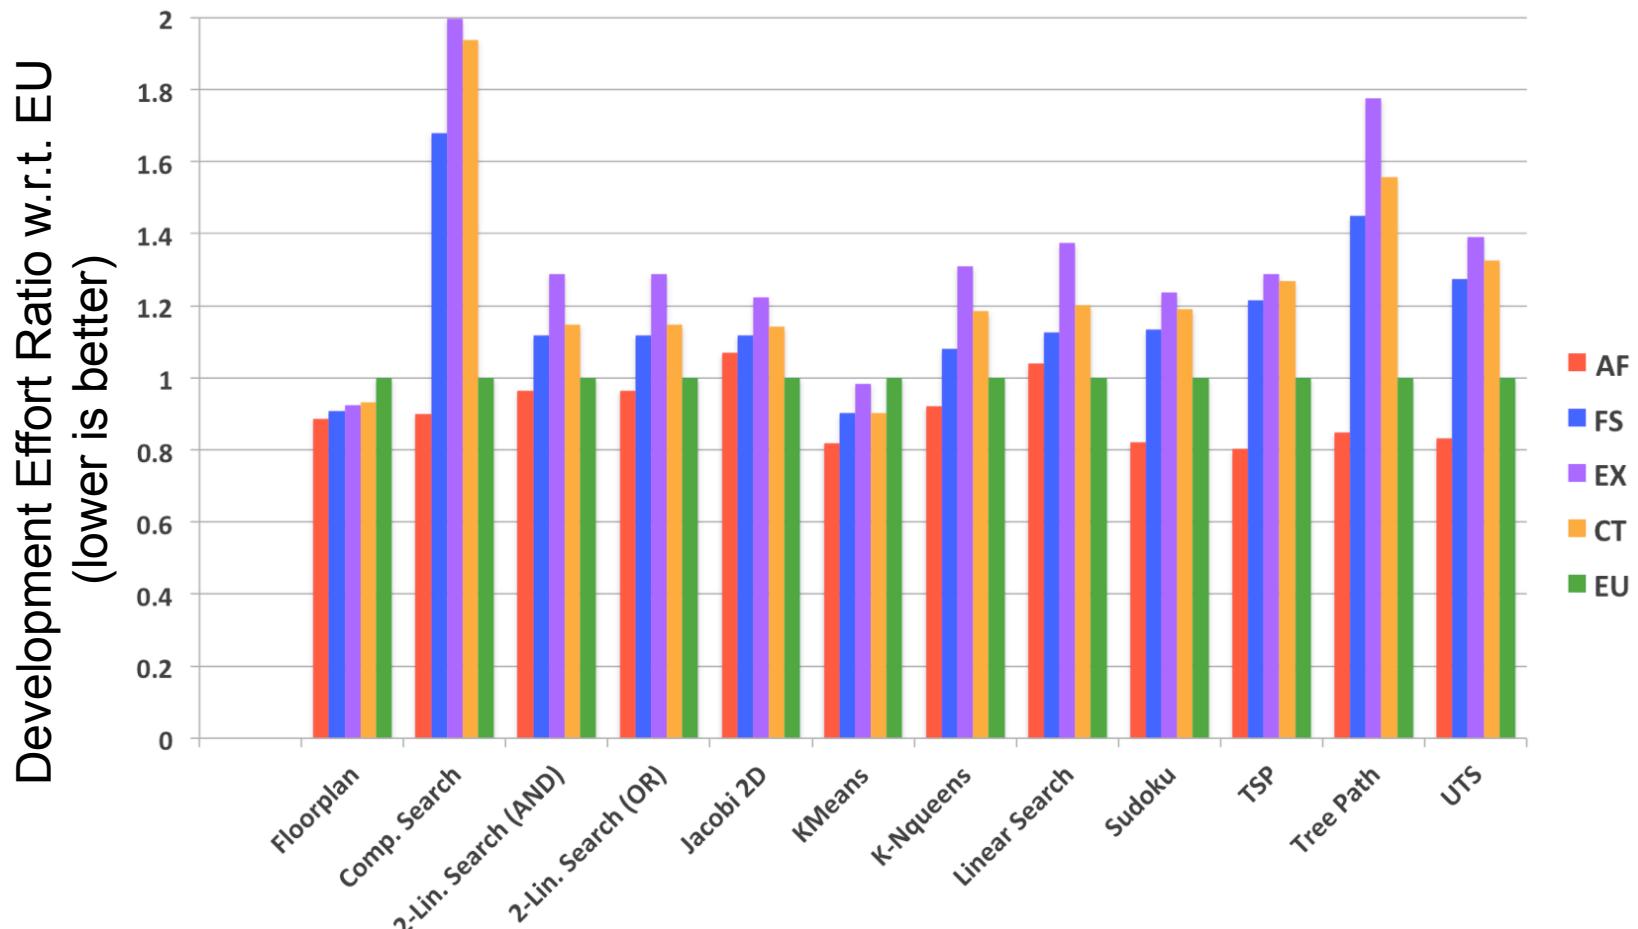
- EU slightly faster than EX and CT, about 50% faster than FS

Productivity – Cyclomatic Complexity



- EU similar to AF, at least 10% better than FS, EX, and CT

Productivity – Halstead's Development Effort



- EU 10% more than AF, 15% lower than FS, EX, and CT.

Eureka Contributions

- **Generalized Eureka-Style Computations**
 - Identified various Eureka patterns
- Proposed Eureka Programming Model
 - Rely on **cooperative** termination techniques
 - Approach works in the presence of inaccessible functions
- Evaluation shows implementation is **effective**
 - Despite using unmodified VM
- Displayed **productivity** benefits via metrics (CC & DE)

The Eureka Programming Model for Speculative Task Parallelism.

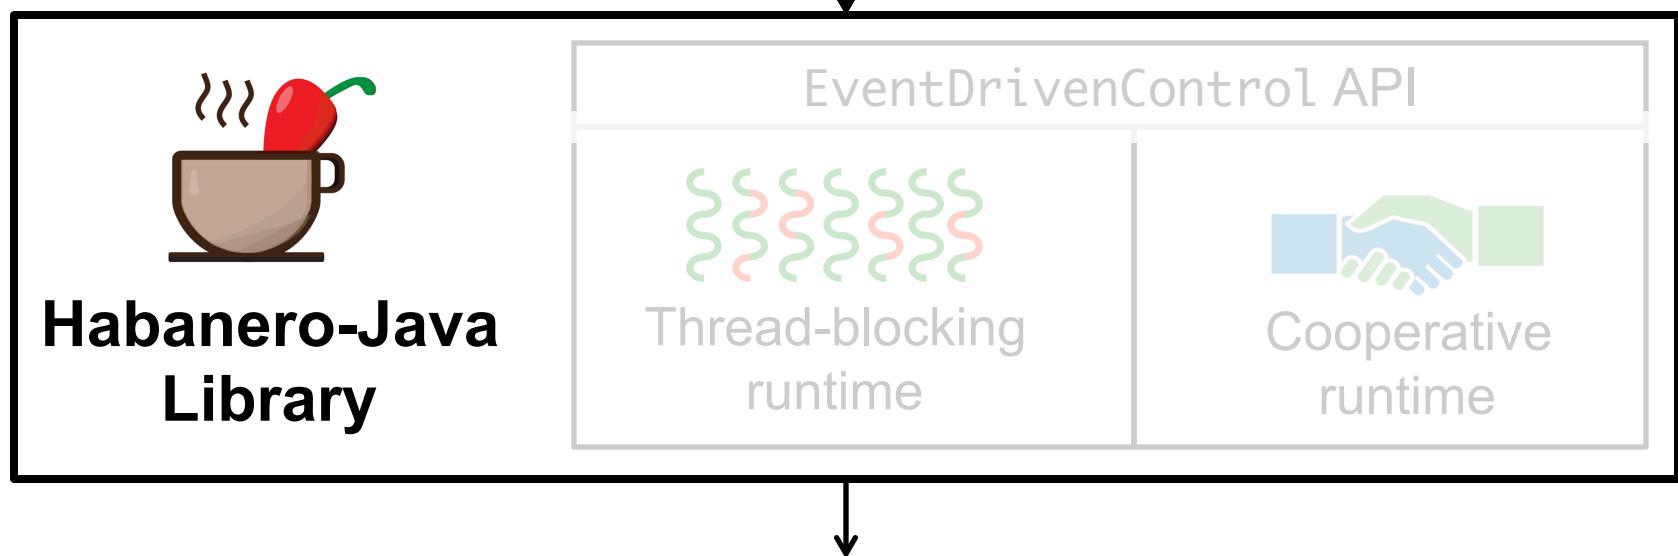
Shams Imam, Vivek Sarkar. 29th European Conference on Object-Oriented Programming (**ECOOP**), July 2015.

Selectors Programming Model

Programming Models

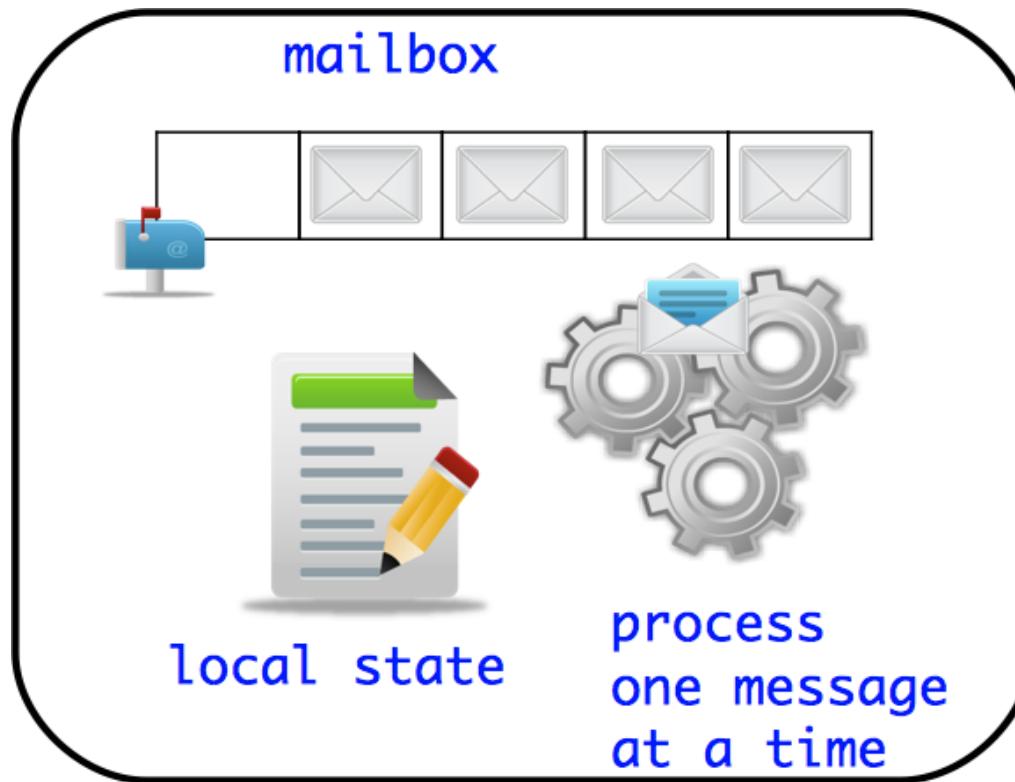


Runtimes



Actors

- Asynchronous message-based concurrency model
- Processes one message at a time





Motivation for Selectors

- Actors give strong guarantees about asynchrony and concurrency, but make synchronization and coordination harder
 - Compared to shared-memory model
 - Coordination patterns involving multiple actors are particularly difficult
 - Actor Model (AM) is not a silver bullet
- Until message is processed, solutions may require the actor to
 - Buffer messages
 - Resend messages to itself and rely on fairness guarantees

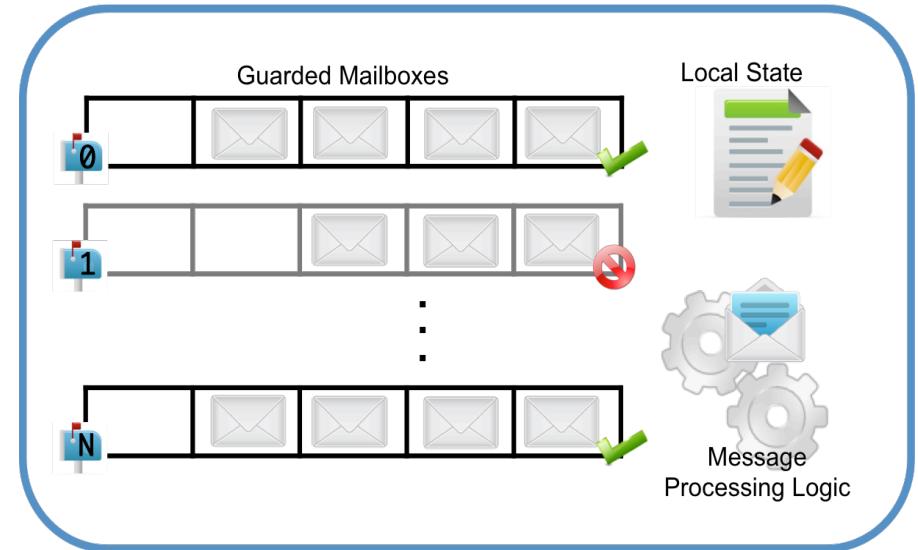
Problematic Synchronization Patterns

- Patterns of interest
 - Join patterns common in streaming applications
 - Synchronous request-reply
 - Priorities in message processing
 - Variants of reader-writer concurrency
 - Producer-consumer with bounded buffer
- Two of these patterns were studied in Master's thesis
 - Relied on different solution of integrating actors and task parallelism

Integrating Task Parallelism with Actors. **Shams Imam**, Vivek Sarkar.
ACM International Conference on Object oriented programming
systems languages and applications (**OOPSLA '12**), October 2012.

Selector

- **Multiple mailboxes**
 - Messages can be concurrently sent to different mailboxes
- Each mailbox maintains a **mutable guard**
 - Mailbox can always receive messages
 - Internal guard changed **cooperatively**
 - Affects which mailboxes provides next message to process

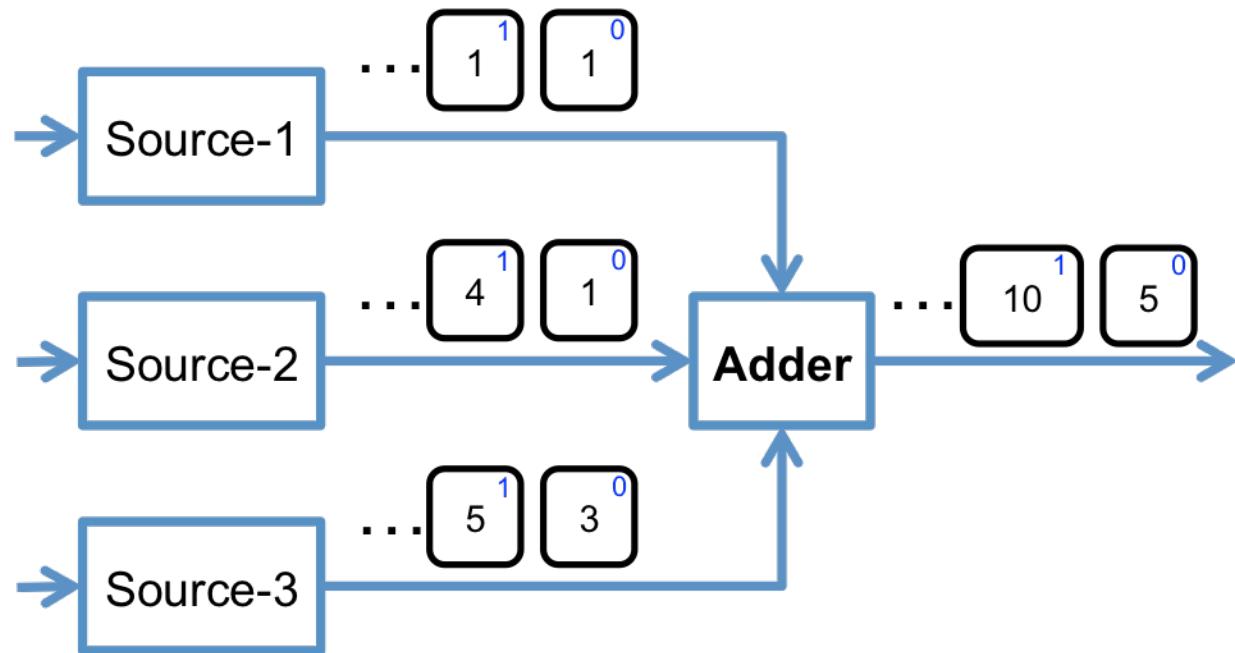


Actor is a Selector with a single mailbox!

- Guard on the mailbox always enabled

Join Patterns in Streaming Applications

- Messages from two or more data streams are combined together into a single message
- Joins need to match inputs from each source
- Wait until all corresponding inputs become available

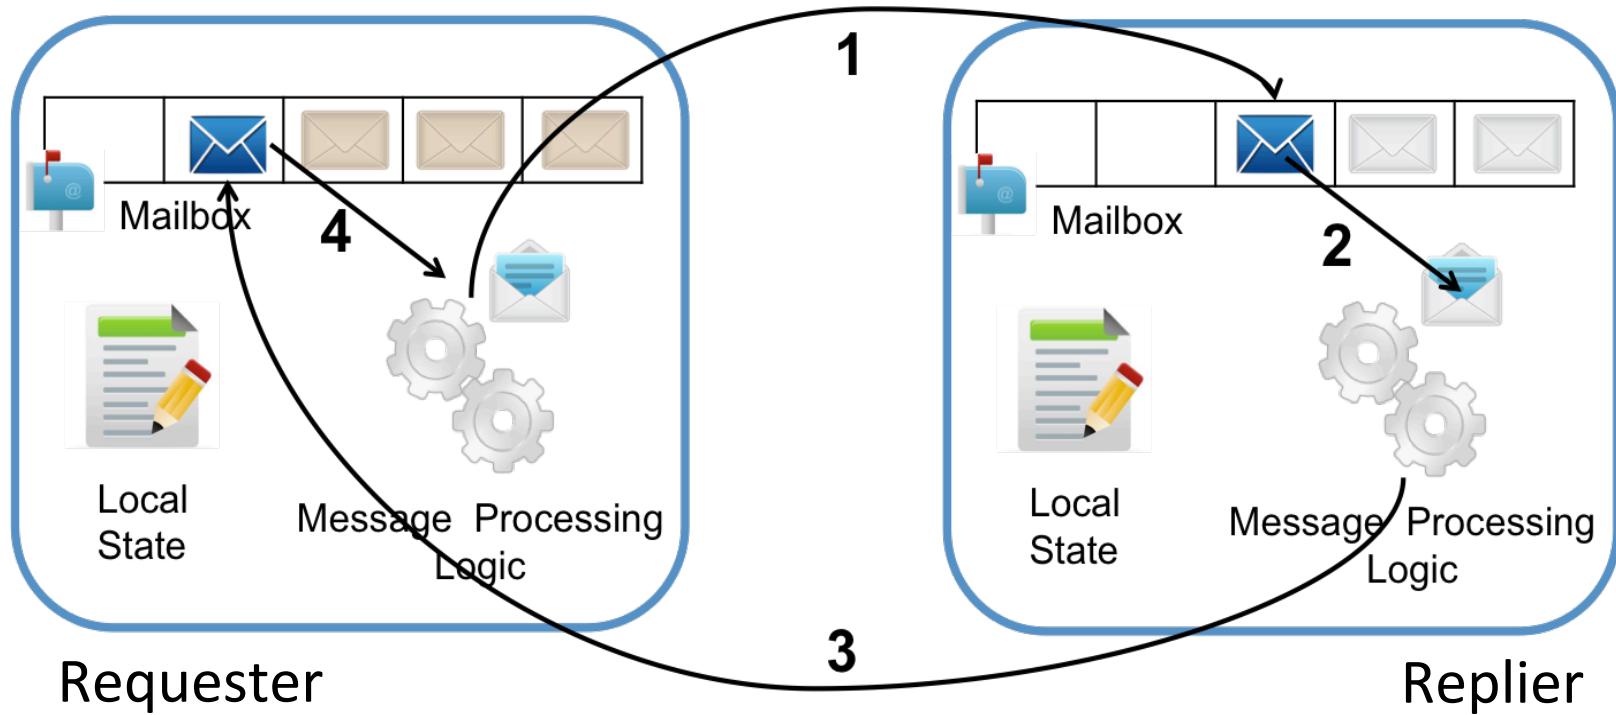


Join Pattern: Selector-based Solution

- One mailbox for each source
- Sources send their messages to corresponding mailboxes
- Two policies
 - Round-Robin order (RR)
 - Receive messages from ordered source for the current sequence
 - Useful if the join actor is performing associative reduction
 - Arbitrary Order (AO)
 - Receive messages from any source for the current sequence
 - Especially useful for non-associative reductions

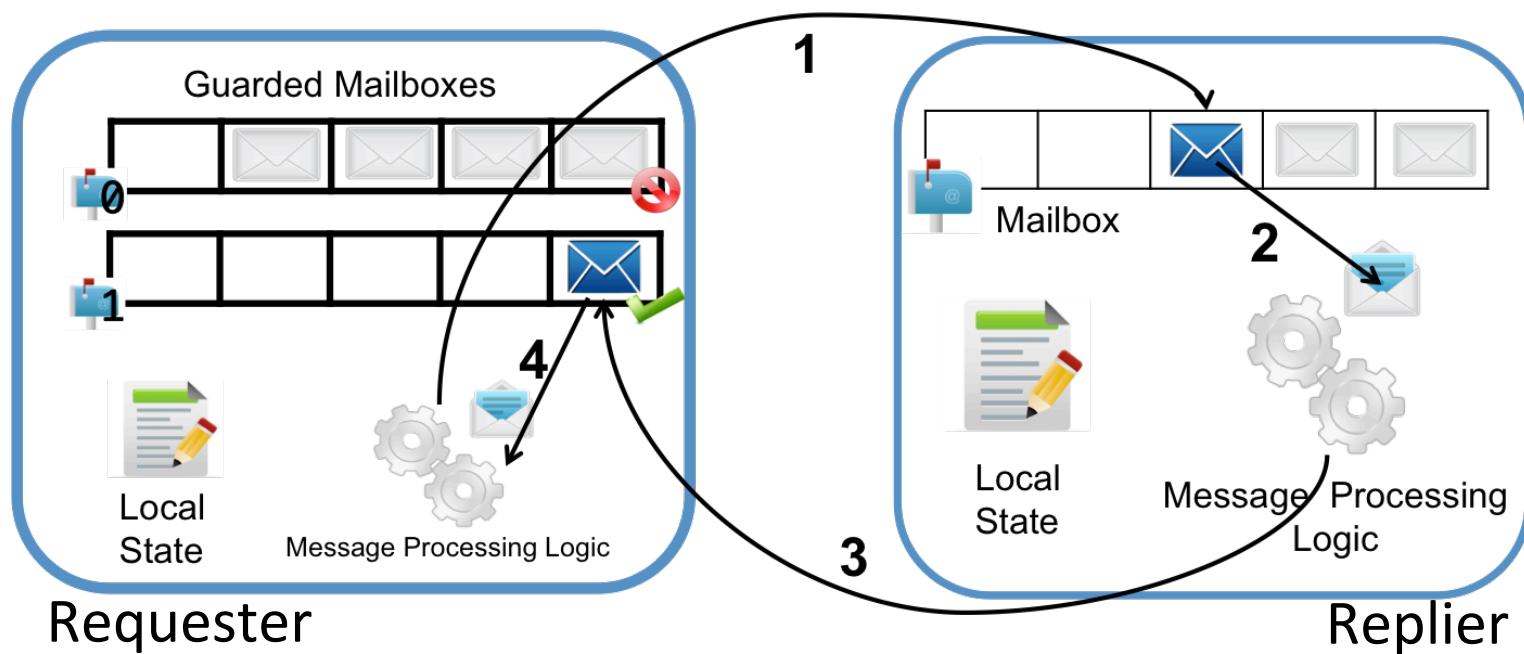
Synchronous Request-Response Pattern

1. Requester sends a message to a replier system
2. Replier *eventually* receives and processes the request
3. Replier returns a message in response
4. Requester can make further progress after receiving response



Synchronous Request/Response: Selector-based Solution

- Two mailboxes
 - one to receive regular messages
 - one to receives synchronous response messages
- Whenever expecting a synchronous response
 - disables the regular mailbox
 - ensures next message processed is from reply mailbox

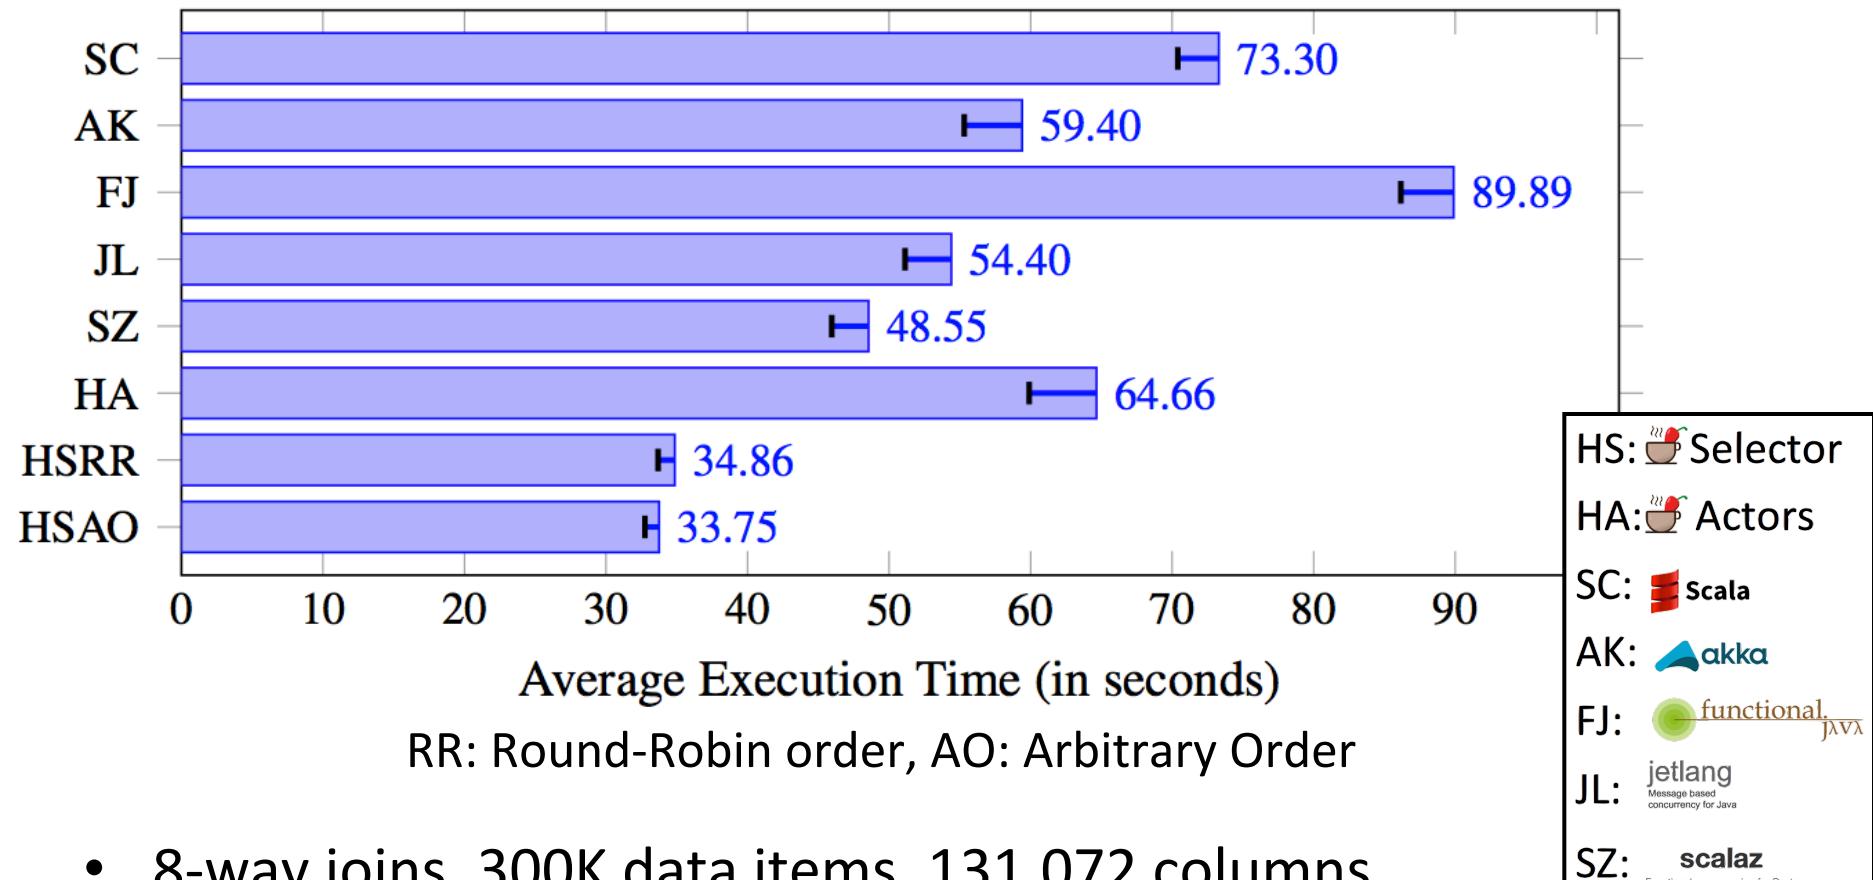


Experimental Results

-  12-core (two hex-cores) 2.8 GHz Intel Westmere SMP node
-  Java Hotspot JDK 1.8.0
- Our implementation:
 -  Selector (HS), pure library implementation on Java 8
- Other libraries:
 -  Actors (HA) 0.1.2
 -  Scala 2.11.0 actors (SC)
 -  akka 2.3.2 (AK)
 -  functional_{java} 4.1 (FJ)
 -  jetlang 0.2.12 (JL)
 -  scalaz 7.1.0-M6 (SZ)
Functional programming for Scala
- Benchmarks from:
Savina – An Actor Benchmark Suite. **Shams Imam**, Vivek Sarkar. 4th Int'l Workshop on Programming based on Actors, Agents, and Decentralized Control (**AGERE! 2014**), October 2014.

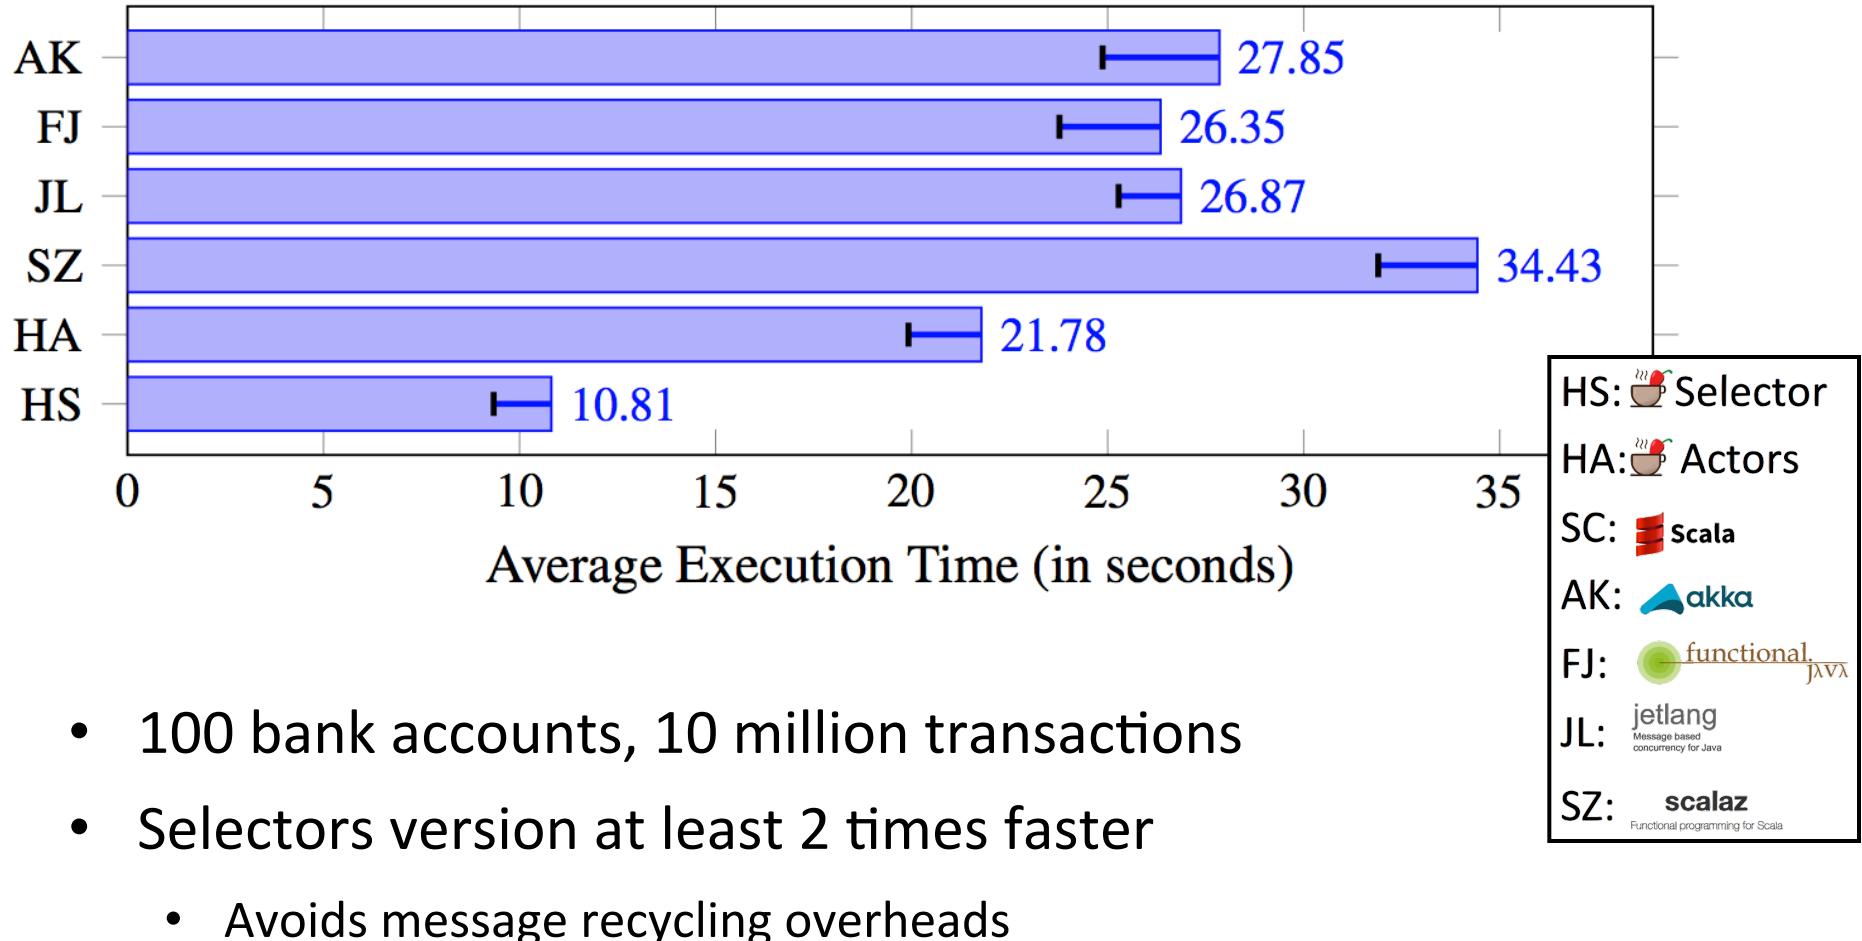


Filter Bank benchmark



- 8-way joins, 300K data items, 131,072 columns
- Selectors version at least 40% faster
 - Avoids local dictionary lookup overheads

Bank Transaction benchmark



Related Work for Selectors

- Pattern matching on receive
 - Enabled-sets by Tomlinson and Singh [OOPSLA '89]
 - Scala Actors by Haller [COORDINATION 2007]
- Aggregator Pattern from Akka
 - Does not match sender
- Message priorities
 - SALSA provides two-level priority [ACM SIGPLAN Notices, Dec 2001]
- Parallel Actor Monitors [Science of Comp. Prog., Feb 2014]
 - Solves the symmetric reader-writer problems
 - Does not support priorities, hence other variants

Selectors Contributions

- Simplify writing of synchronization and coordination patterns
 - Using a simple extension to Actors
 - Multiple cooperatively guarded mailboxes
- Effective solution to synchronization/coordination patterns
 - Join patterns common in streaming applications
 - Synchronous request-reply
 - Variants of reader-writer concurrency
 - Priorities in message processing
 - Producer-consumer with bounded buffer

Selectors: Actors with Multiple Guarded Mailboxes. Shams Imam,
Vivek Sarkar. 4th Int'l Workshop on Programming based on Actors,
Agents, and Decentralized Control (**AGERE! 2014**), October 2014.



Thesis Summary

- **Cooperative techniques** help parallel programs
 - In presence of modern **synchronization and coordination constructs**
- **Cooperative Runtime**
 - Avoiding thread-blocking operations
 - Enable implementing future synchronization constructs
- Programming Model **Extension: Eureka**
 - Efficient collective cooperative task termination
 - Performance by reducing redundant work
 - Productivity by simplifying user-level code
- Programming Model **Extension: Selector**
 - Performance by reducing book-keeping performed by user
 - Productivity by simplifying user-level code

Acknowledgments

- Vivek Sarkar
- Other Committee members
- Habanero Research Group
- Colleagues at Rice CS Department

Questions

- Cooperative techniques help parallel programs
 - In presence of modern synchronization and coordination constructs
- Cooperative Runtime

import defense.audience.Questions;
• Avoiding thread-blocking operations
• Enable implementing future synchronization constructs

import defense.audience.Comments;

- Efficient collective cooperative task termination
- Performance by reducing redundant work
- Productivity by simplifying user-level code
- Programming Model Extension: Selector
 - Performance by reducing book-keeping performed by user
 - Productivity by simplifying user-level code